

# De ARIMA aux Transformers

## Analyse de Séries Temporelles et NLP Avancé

Réalisée par :

- KHATTABI Mohammed
- OUAHAB Abderrahim
- ZOURDAN Youssef

2025/03/07

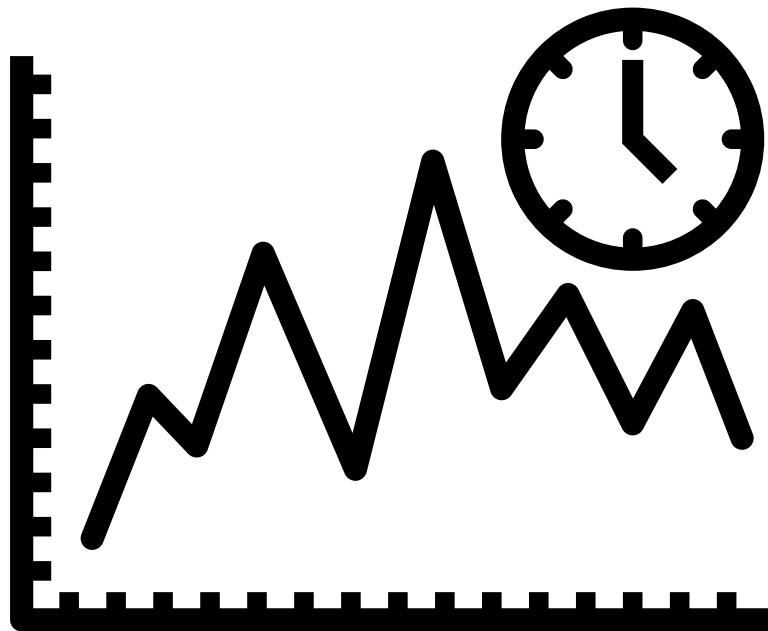
Encadré par :

- M. Abdelwahab Naji

# Sommaire

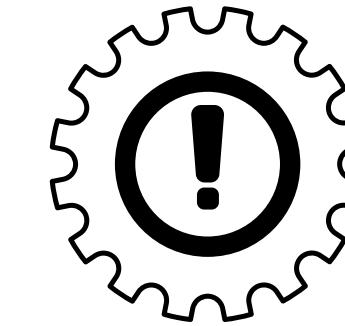
- 1. Time Series : Objectifs et Défis**
- 2. Premiers Pas en Time Series**
- 3. Découverte des Jeux de Données**
- 4. Deep Learning : ANN, RNN, et LSTM**
- 5. Analyse de Séries Temporelles et Modèles Avancés**
- 6. NLP et Techniques Avancées**
- 7. Perspectives et Conclusions**

# 1. Time Series : Objectifs et Défis



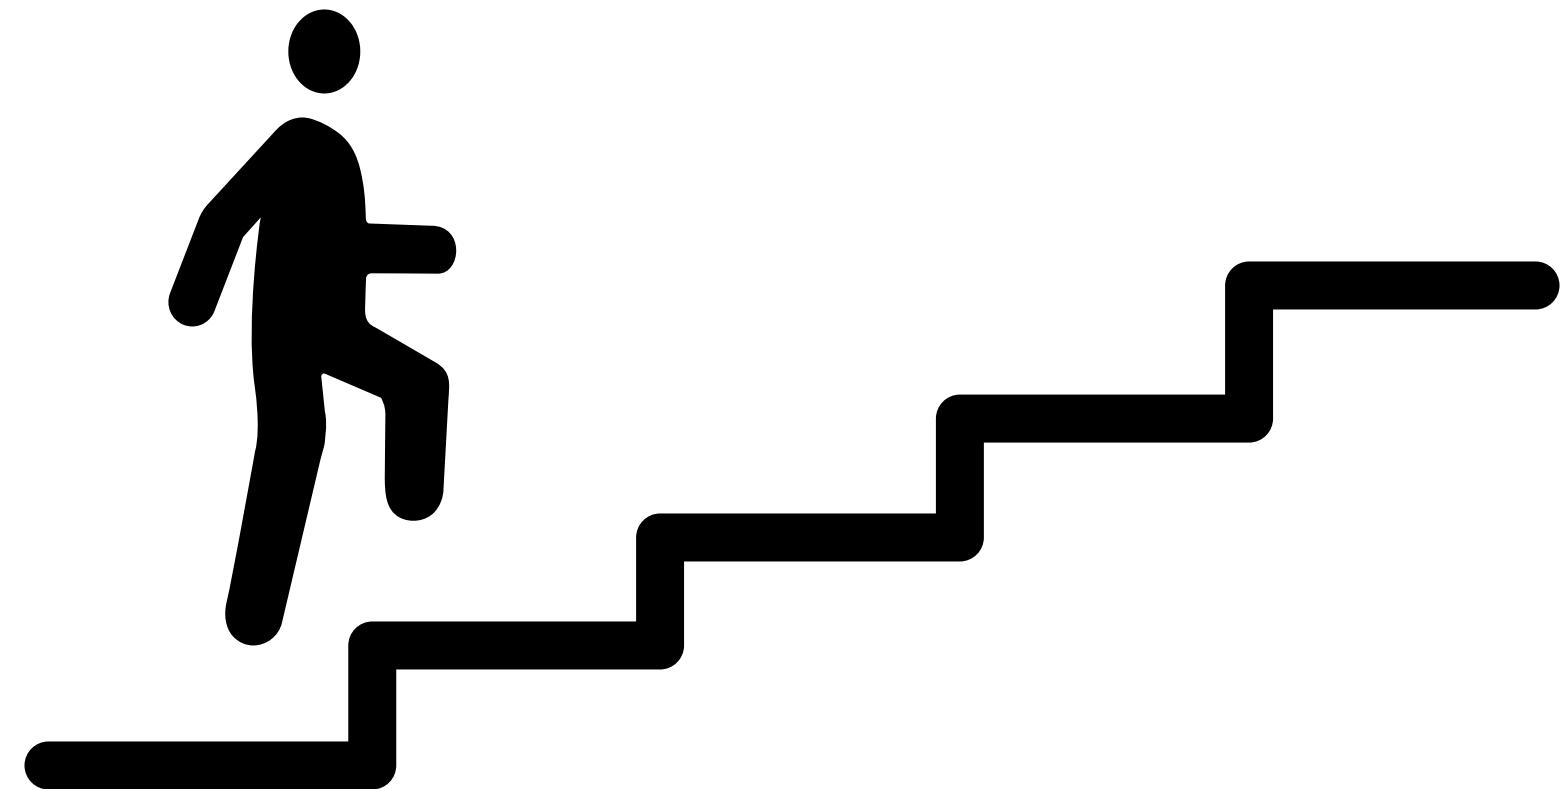


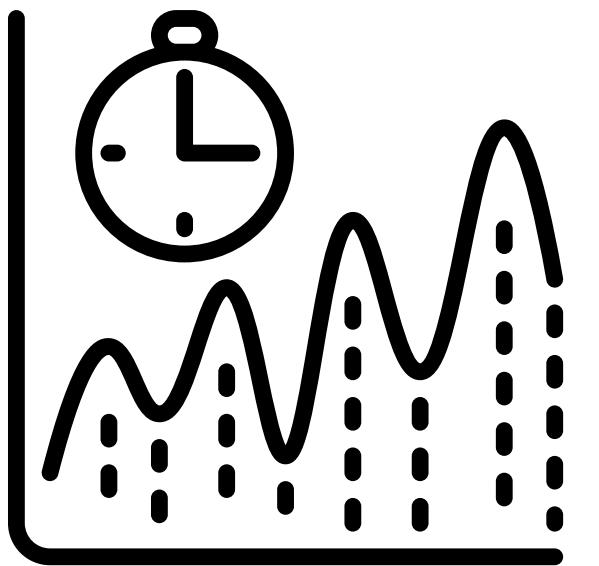
- Compréhension des séries temporelles
- Préparation des données
- Visualisation et analyse
- Modélisation et prévision
- Optimisation et évaluation



- Dépendance temporelle et identification des tendances et saisonsnalités.
- Stationnarité et transformations pour stabiliser la série.
- Choix du bon modèle
- Gestion des erreurs résiduelles et surajustement.
- Problèmes liés aux RNN/LSTM : vanishing gradient, overfitting, convergence.

## 2. Premiers Pas en Time Series



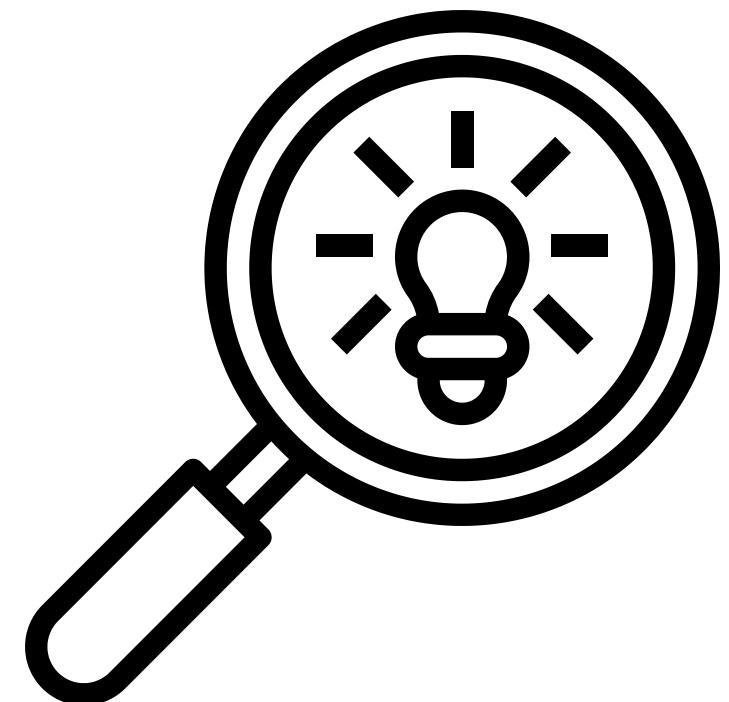


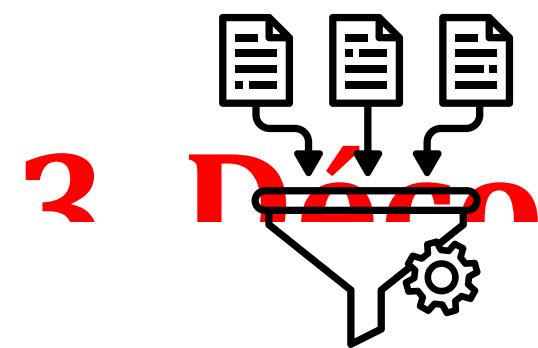
**Données organisées chronologiquement.**

- Collecte et visualisation des données.
- Identification des tendances et saisonnalités.
- Modélisation et prévision des données.

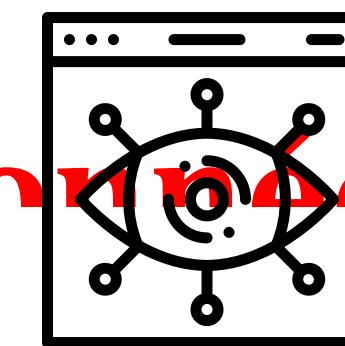
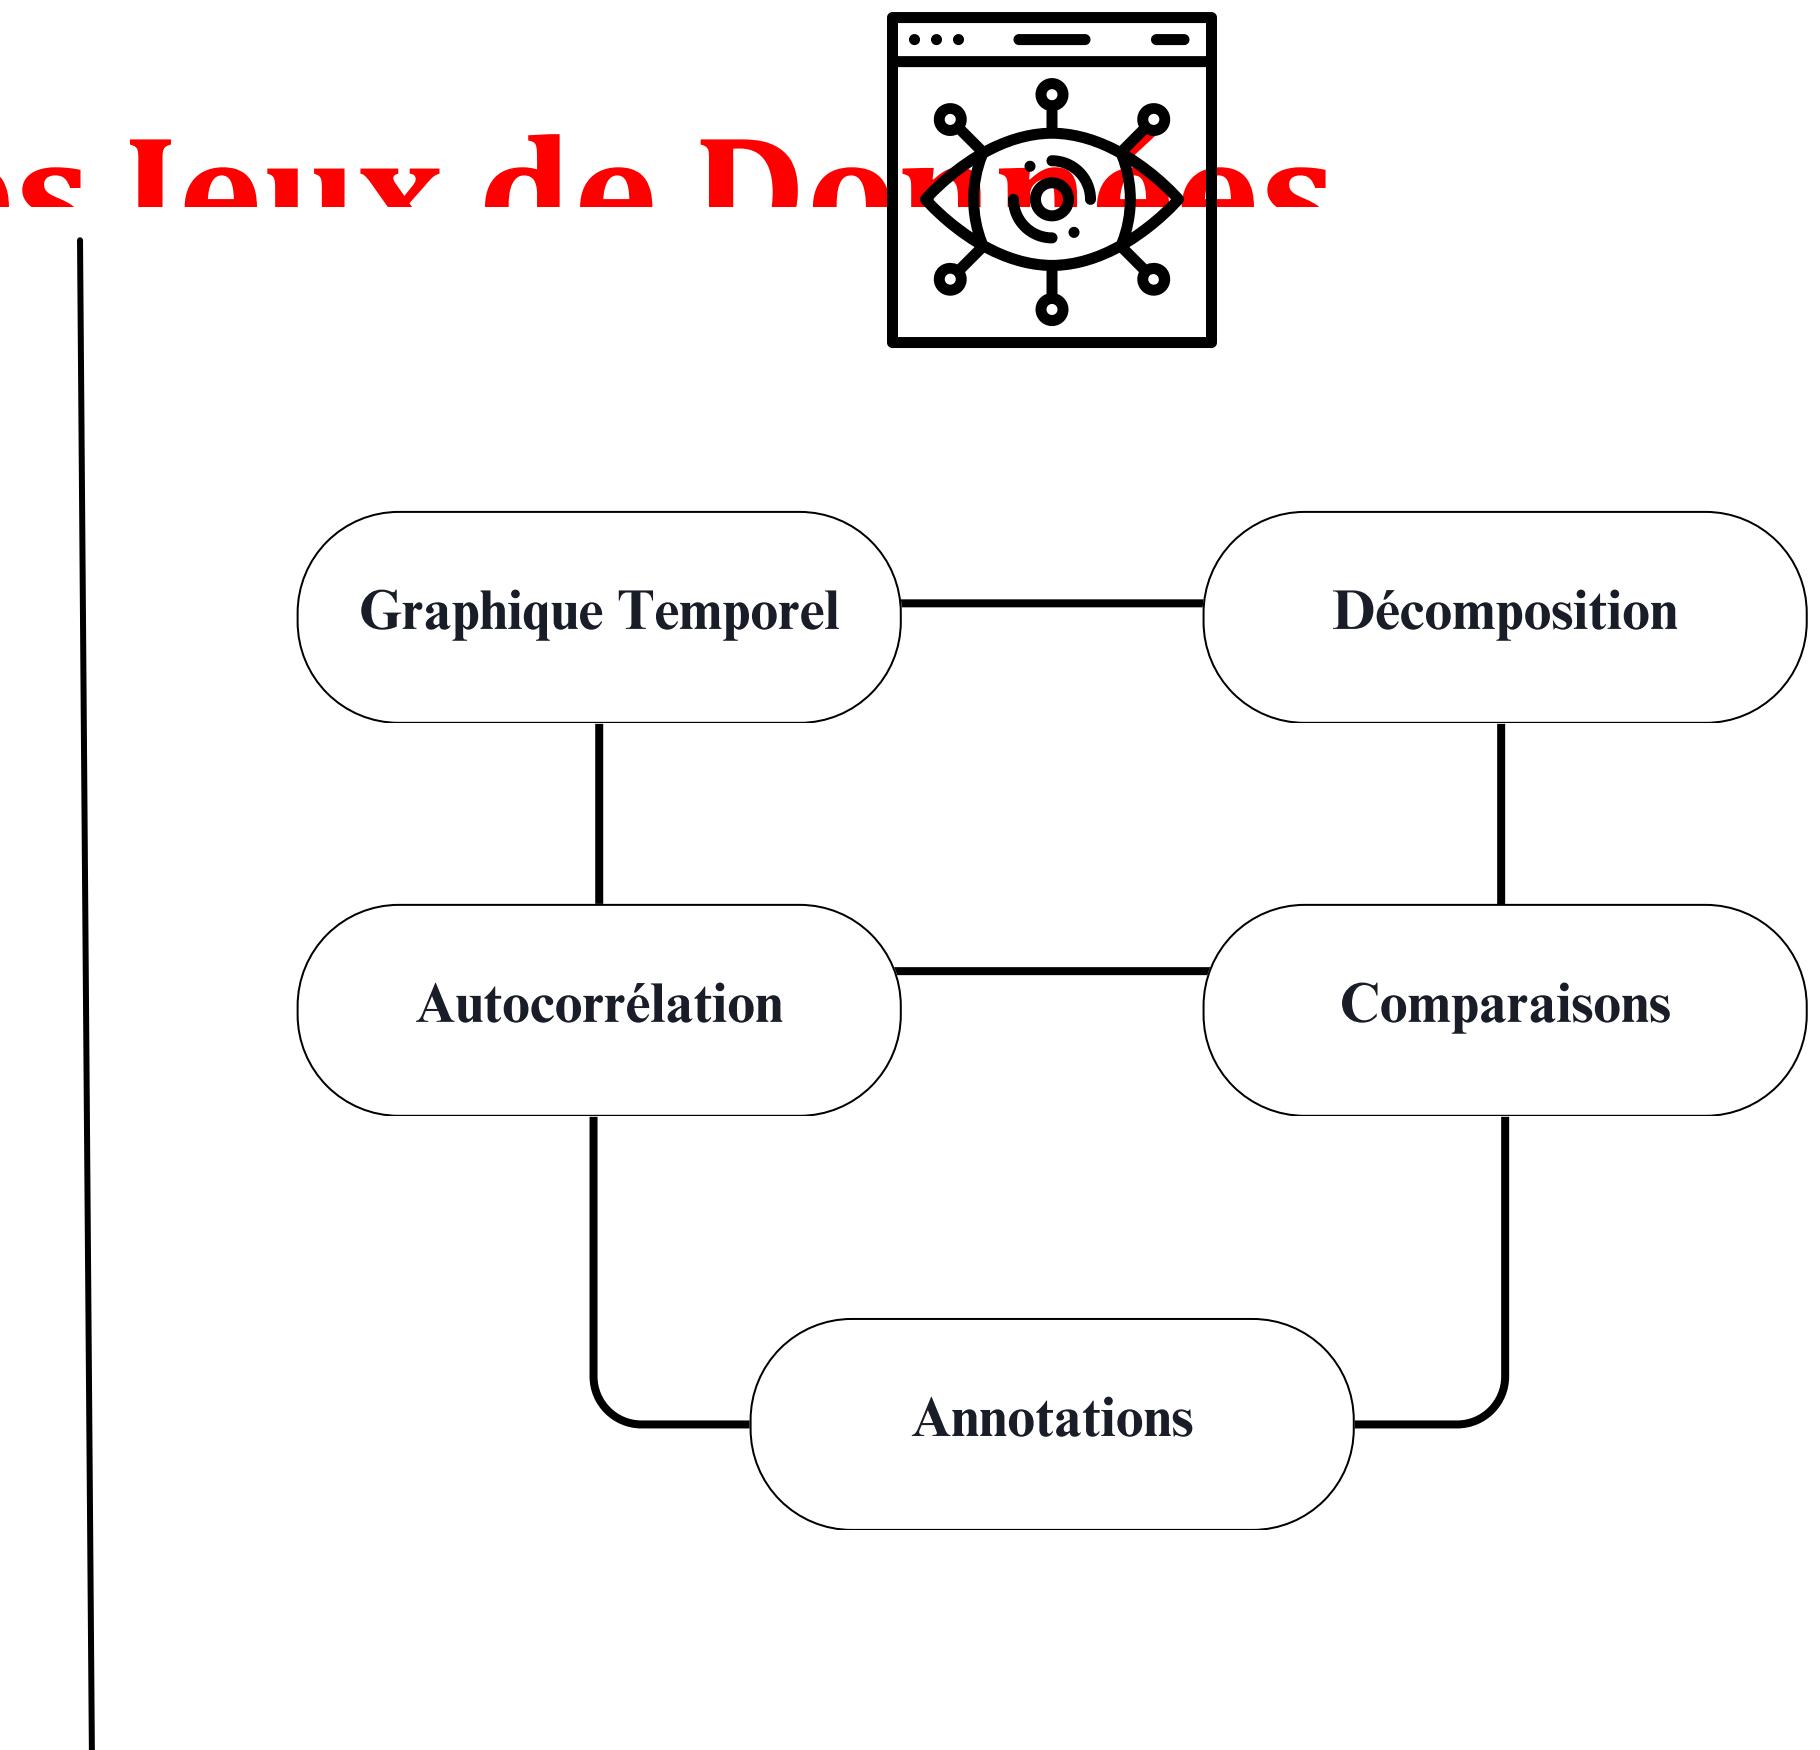
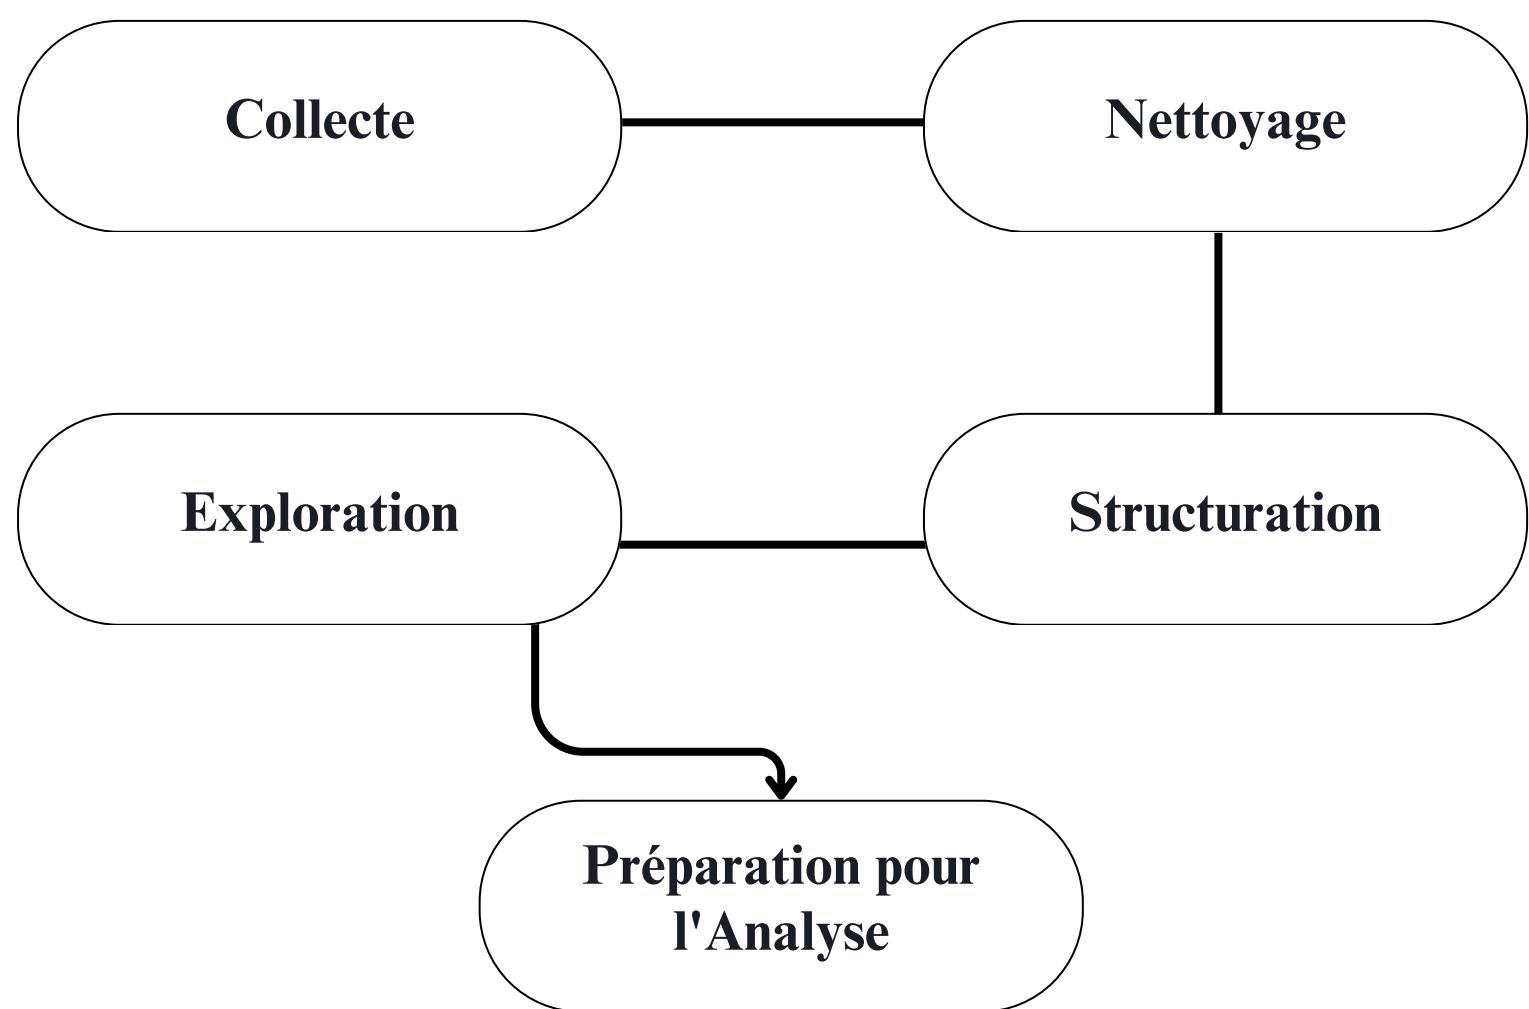


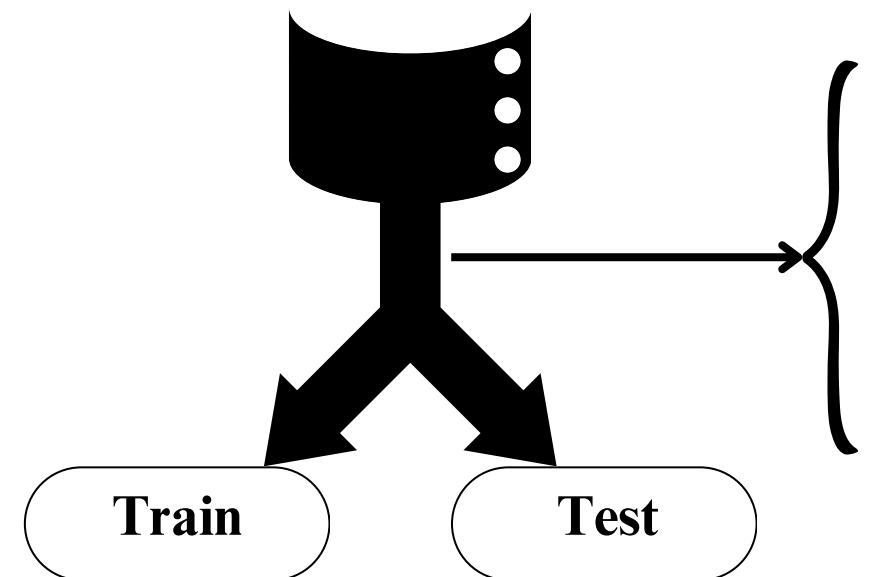
### **3. Découverte des Jeux de Données**





## 3 Découverte des Liens de Données





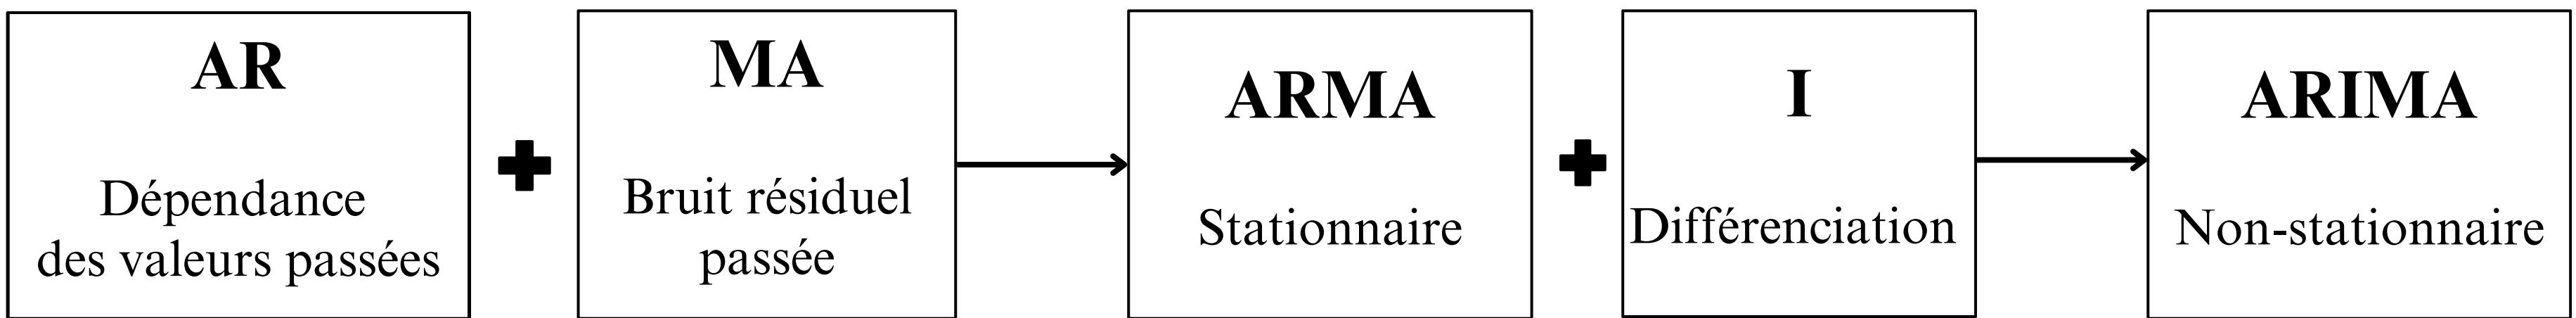
- Séparation temporelle
- Ratio classique

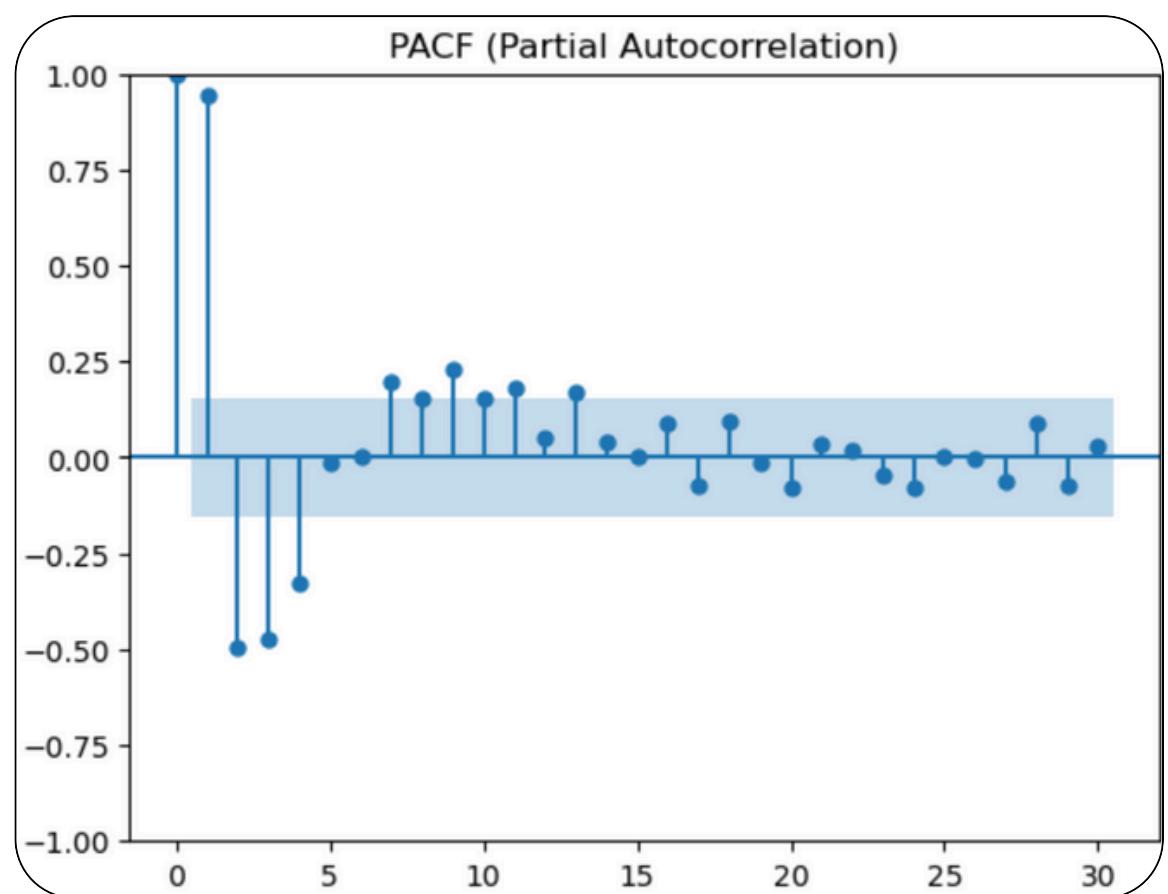
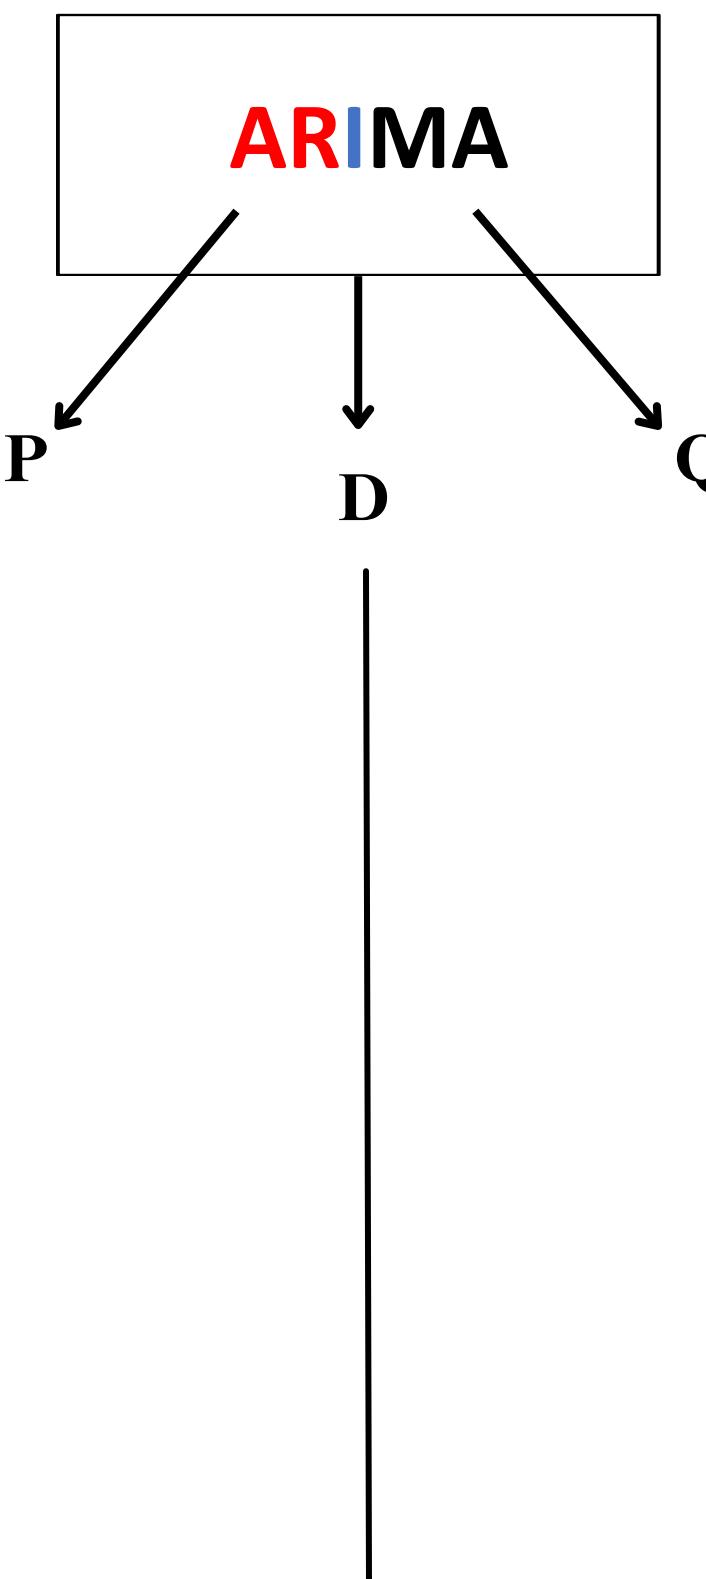


- Eviter les fuites de données
- Taille du test

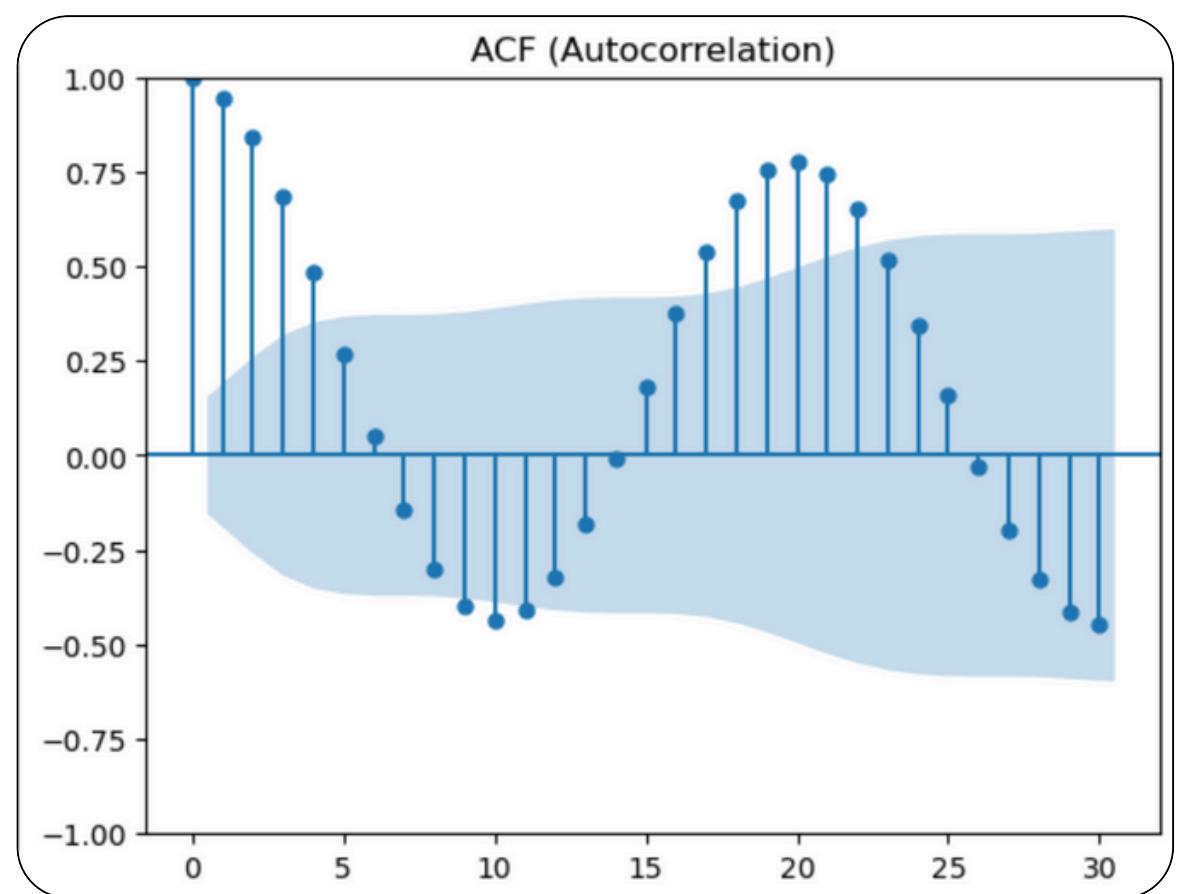
```
# Charger les données
data = pd.read_csv('data.csv', parse_dates=['date'], index_col='date')

# Diviser les données
train_size = int(len(data) * 0.8) # 80% pour l'entraînement
train, test = data[:train_size], data[train_size:]
```





corrélation directe entre une série et un lag spécifique

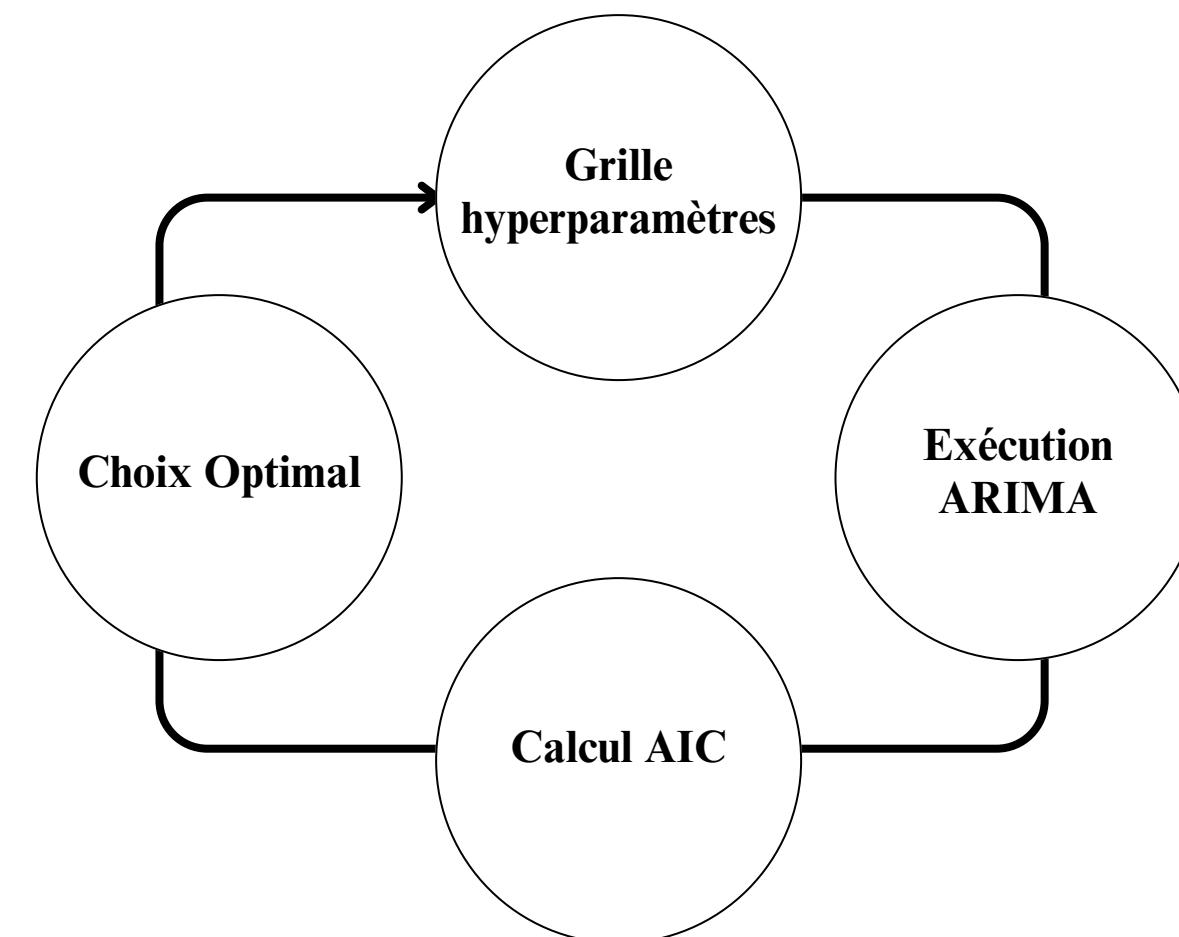


corrélation entre une série temporelle et ses lags

# Grid Search

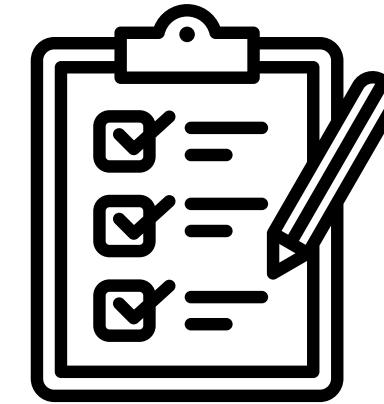


Trouver les meilleures valeurs des hyperparamètres



```
p=2, d=0, q=0 -> AIC=937.2276302268851
...
p=3, d=2, q=1 -> AIC=912.3296883733982
p=3, d=2, q=2 -> AIC=915.533357524179
p=3, d=2, q=3 -> AIC=911.0013625037545
Meilleurs paramètres ARIMA: (0, 2, 3) avec AIC: 907.1687054410694
```

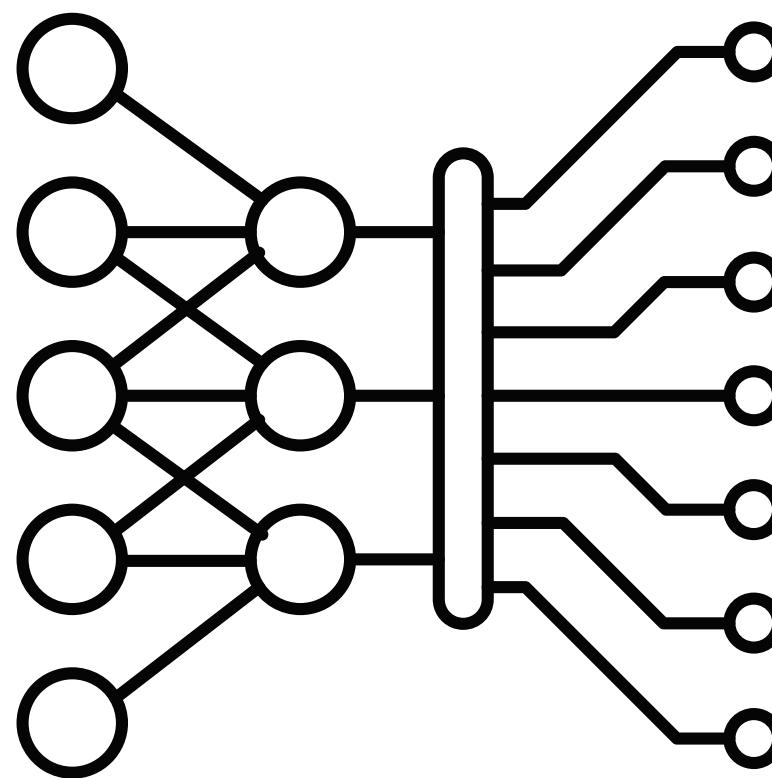
## Évaluation des modèles

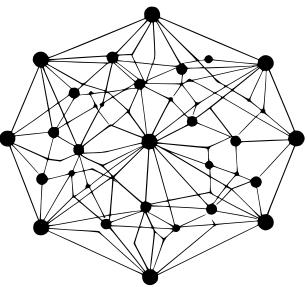


mesurer leur performance et d'identifier les améliorations nécessaires

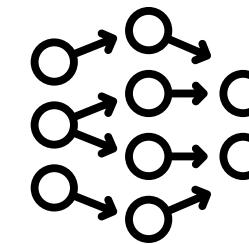
```
Residual Forecast Error: -1.3208072068007628
Forecast Bias: -1.3208072068007628
Mean Absolute Error: 47.09943294486856
Mean Squared Error: 3674.176001779176
Root Mean Squared Error: 60.614981661130415
```

## 4. Deep Learning : ANN, RNN et LSTM

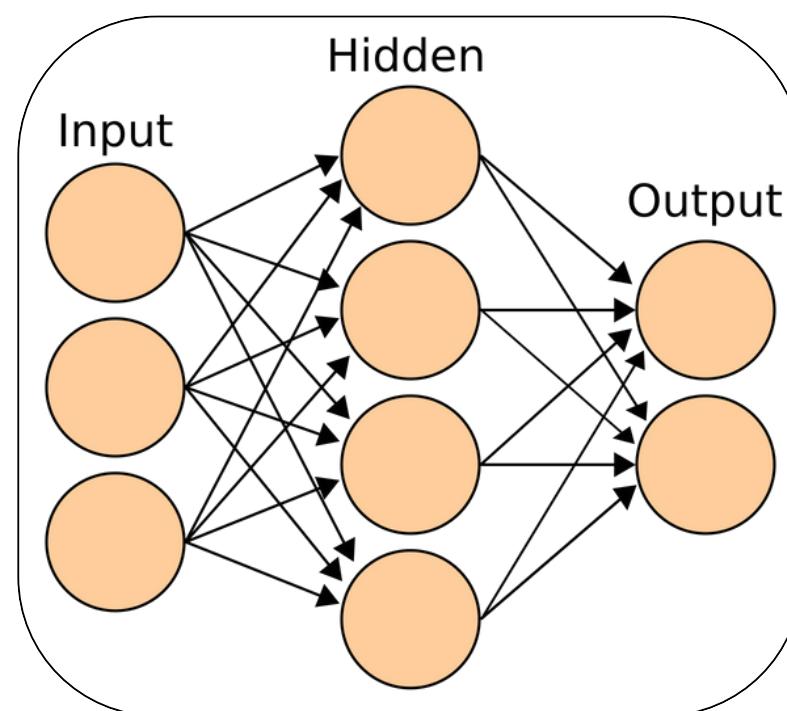




## Deep Learning

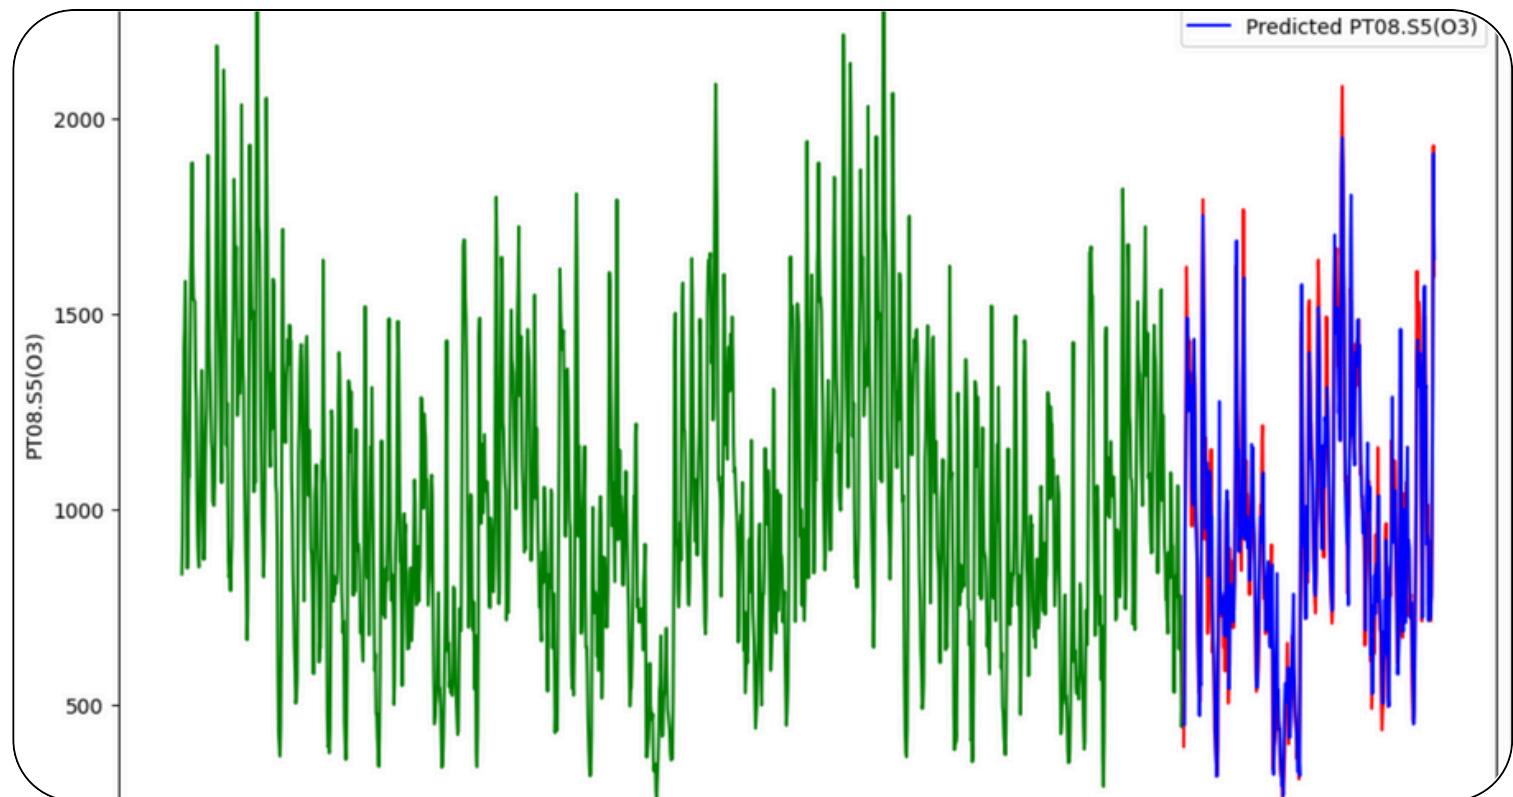
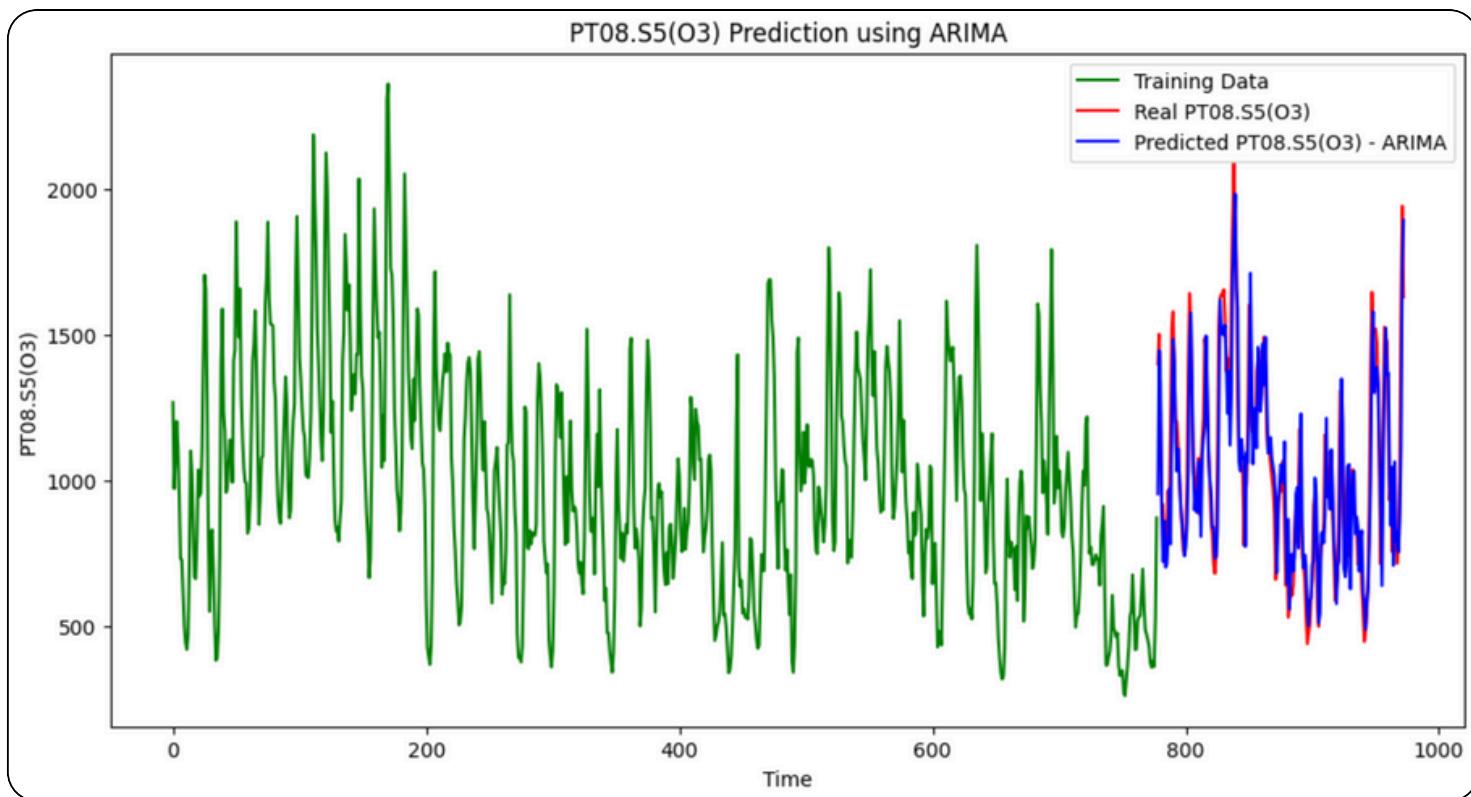


- Le Deep Learning est une approche basée sur les réseaux de neurones pour capturer des relations complexes dans les données.
- Contrairement à ARIMA : il permet de modéliser des tendances non linéaires.



# Avantages

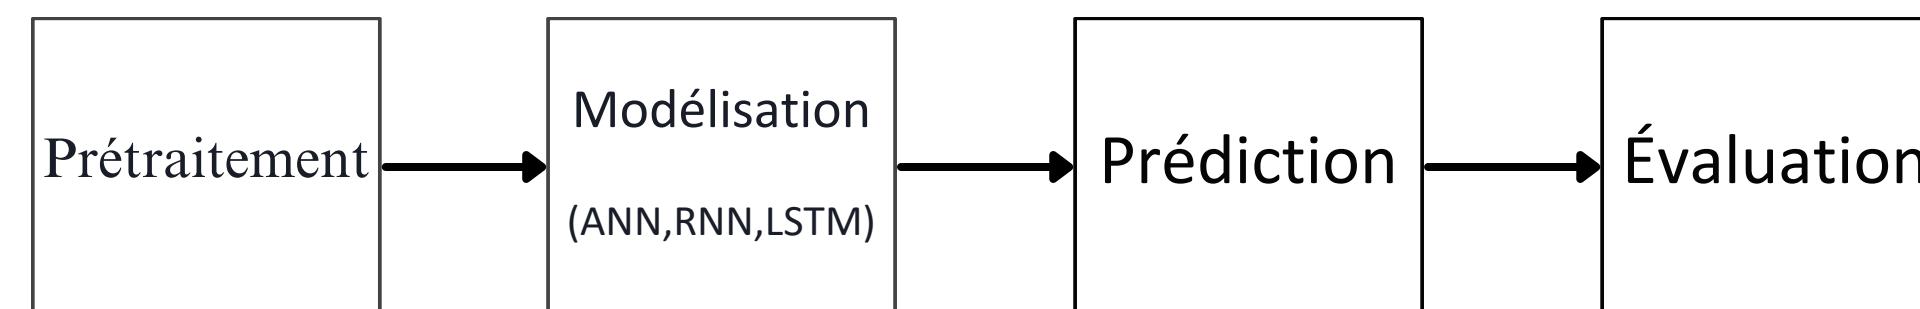
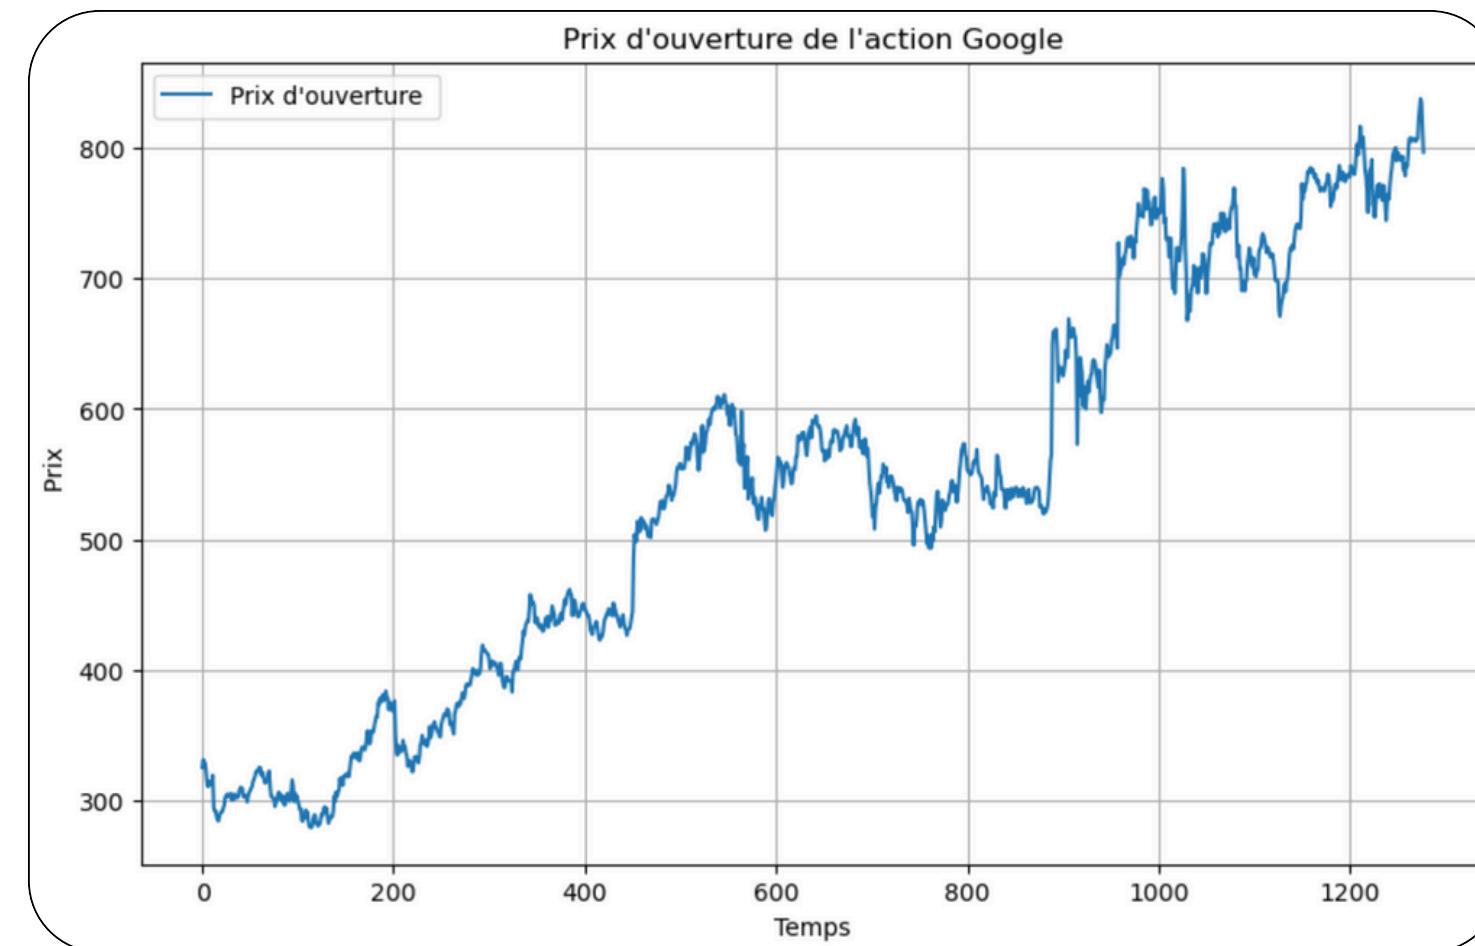
- Capture des dépendances à long terme.
- Gestion des tendances et saisonnalités.



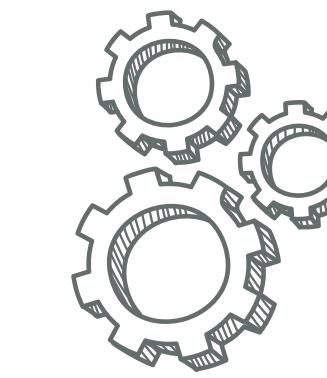
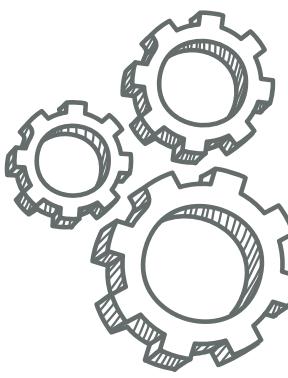
Model	MAE	MSE	RMSE	R2
ARIMA	123.409645	24965.311545	158.004150	0.781957
LSTM	103.505379	17977.183253	134.079019	0.860822

# Google Stock Analysis

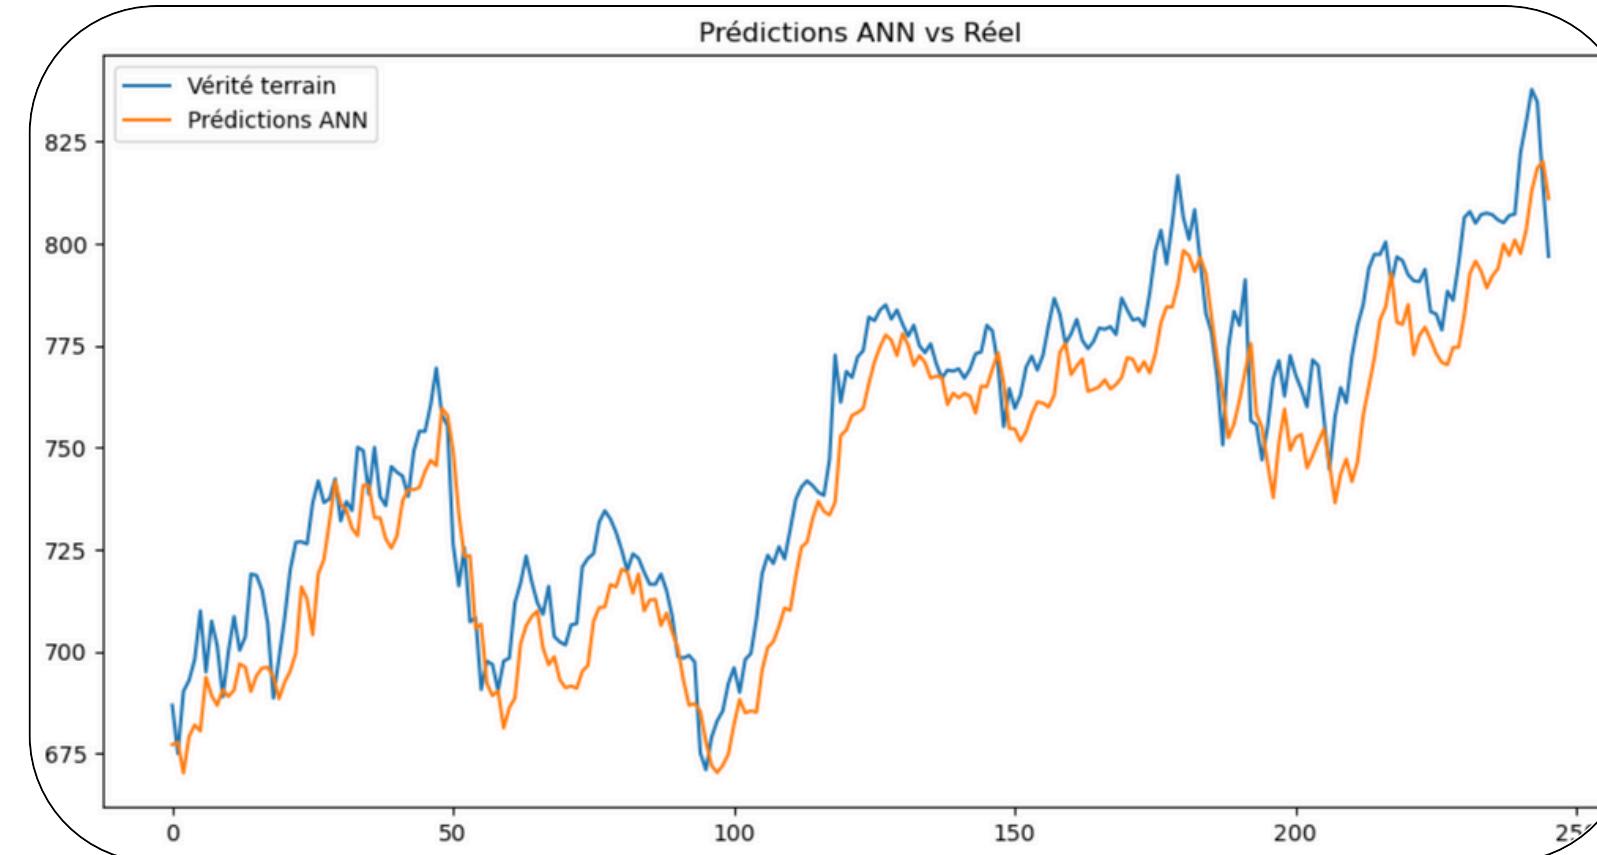
- **Objectif** : Prédire l'évolution du cours des actions Google en utilisant ANN, RNN et LSTM.
- **Données utilisées** : Historique des prix des actions Google “Google Stock Price”.



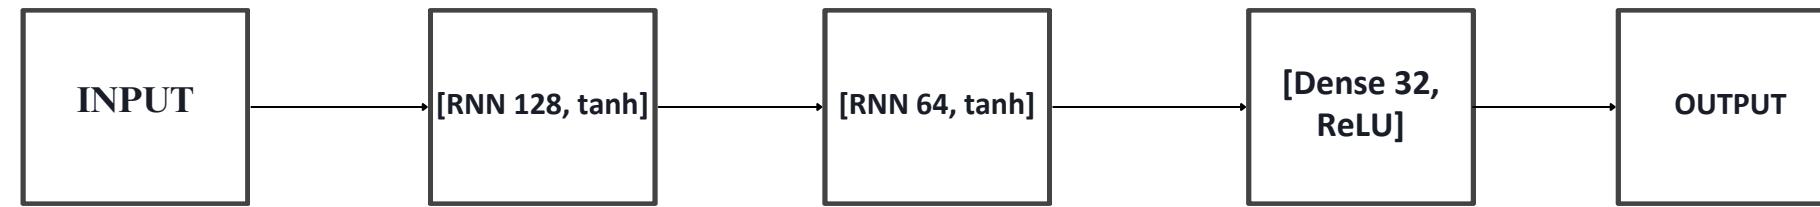
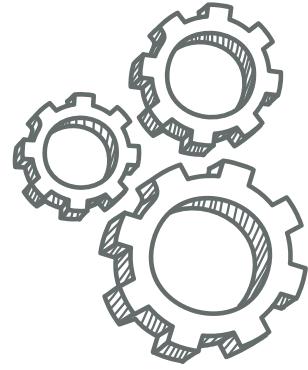
# Modélisation ANN



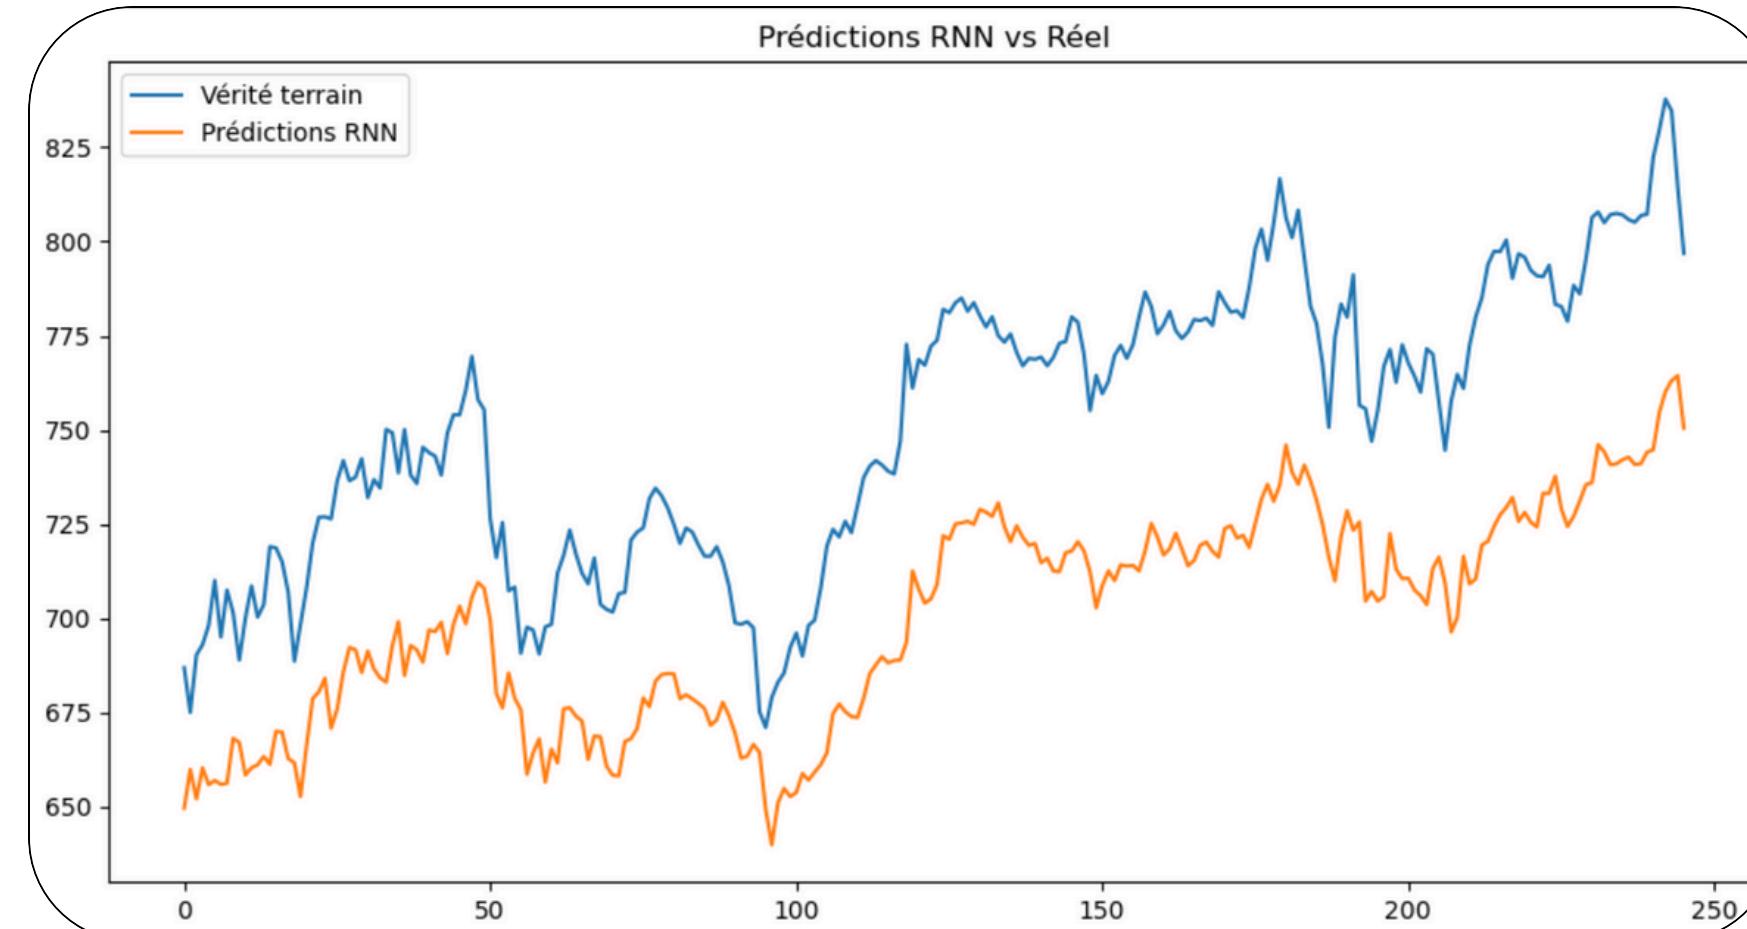
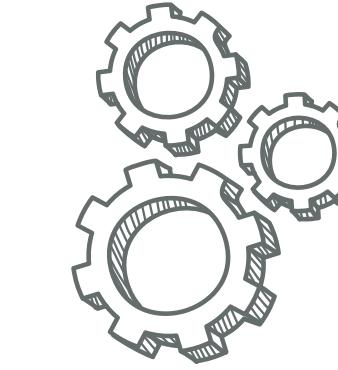
- Optimisation avec Adam et fonction de perte MSE.
- 100 epochs, batch size=32, validation split=20%.



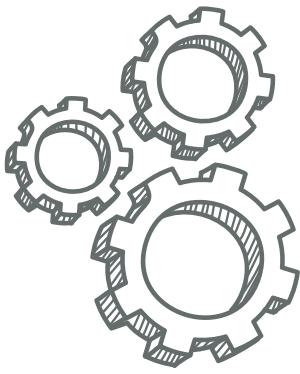
# Modélisation RNN



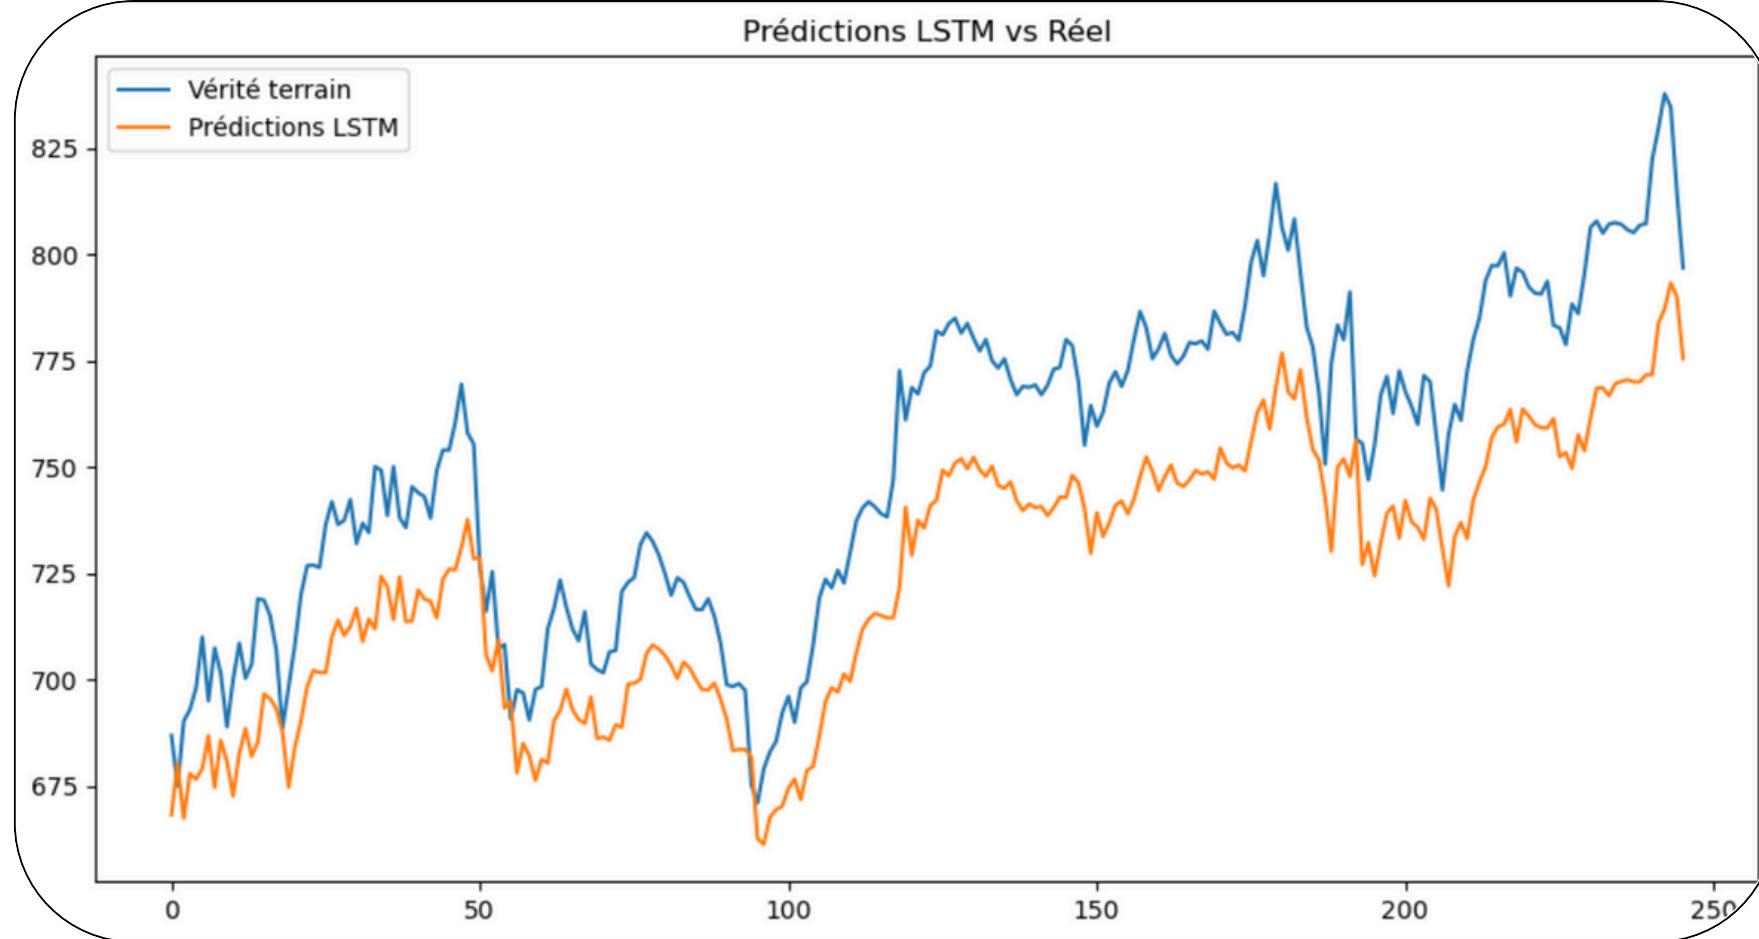
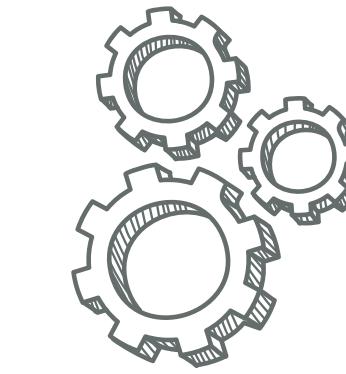
- Optimisation avec Adam et fonction de perte MSE.
- 100 epochs, batch size=32, validation split=20%.



# Modélisation LSTM



- Optimisation avec Adam et fonction de perte MSE.
- 100 epochs, batch size=32, validation split=20%.

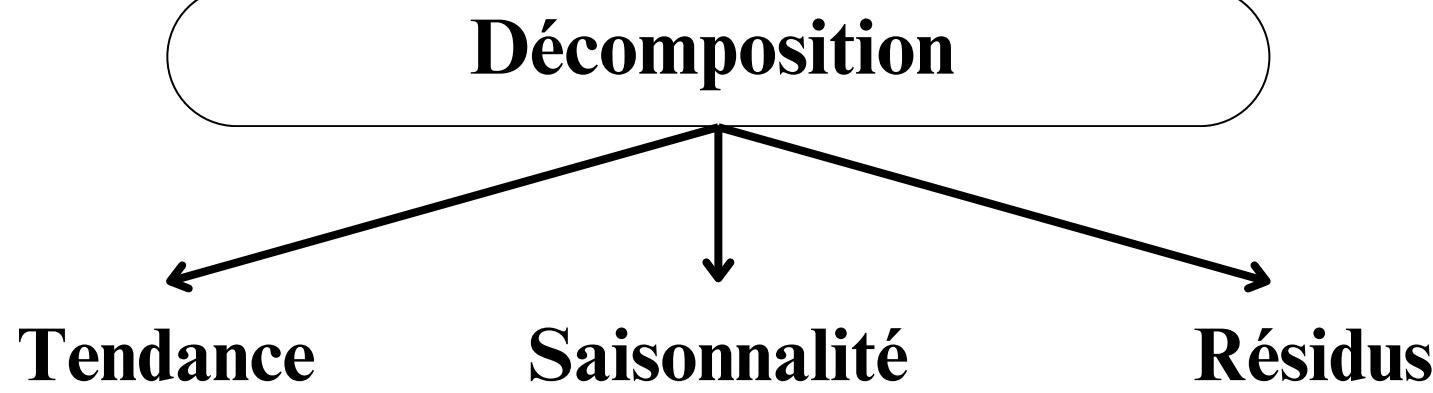


# **5. Analyse de Séries Temporelles et Modèles Avancés**



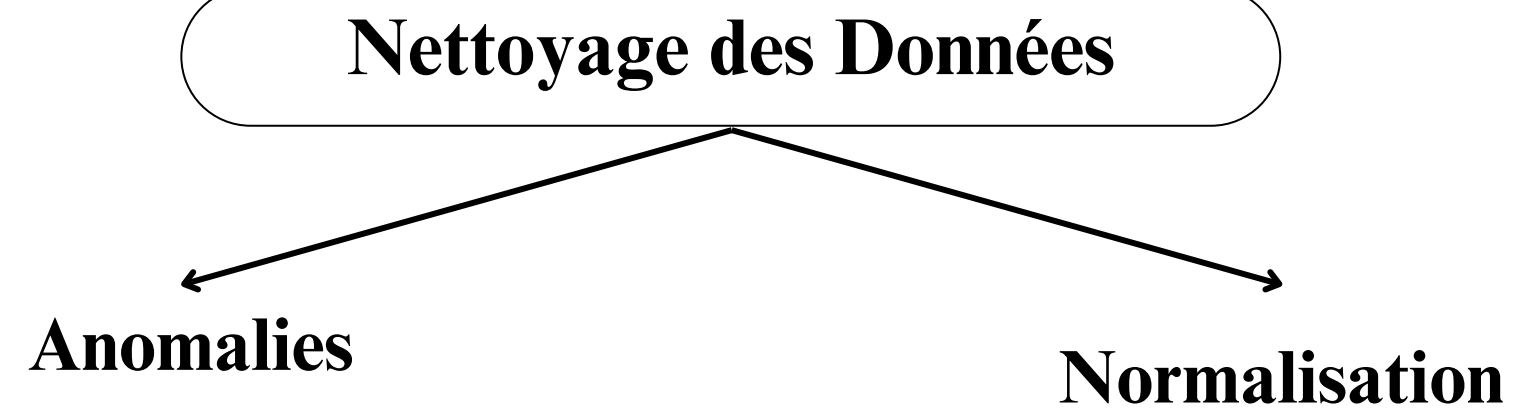
# Présentation du Dataset

Variables clés : Date, PT08.S5(O3)



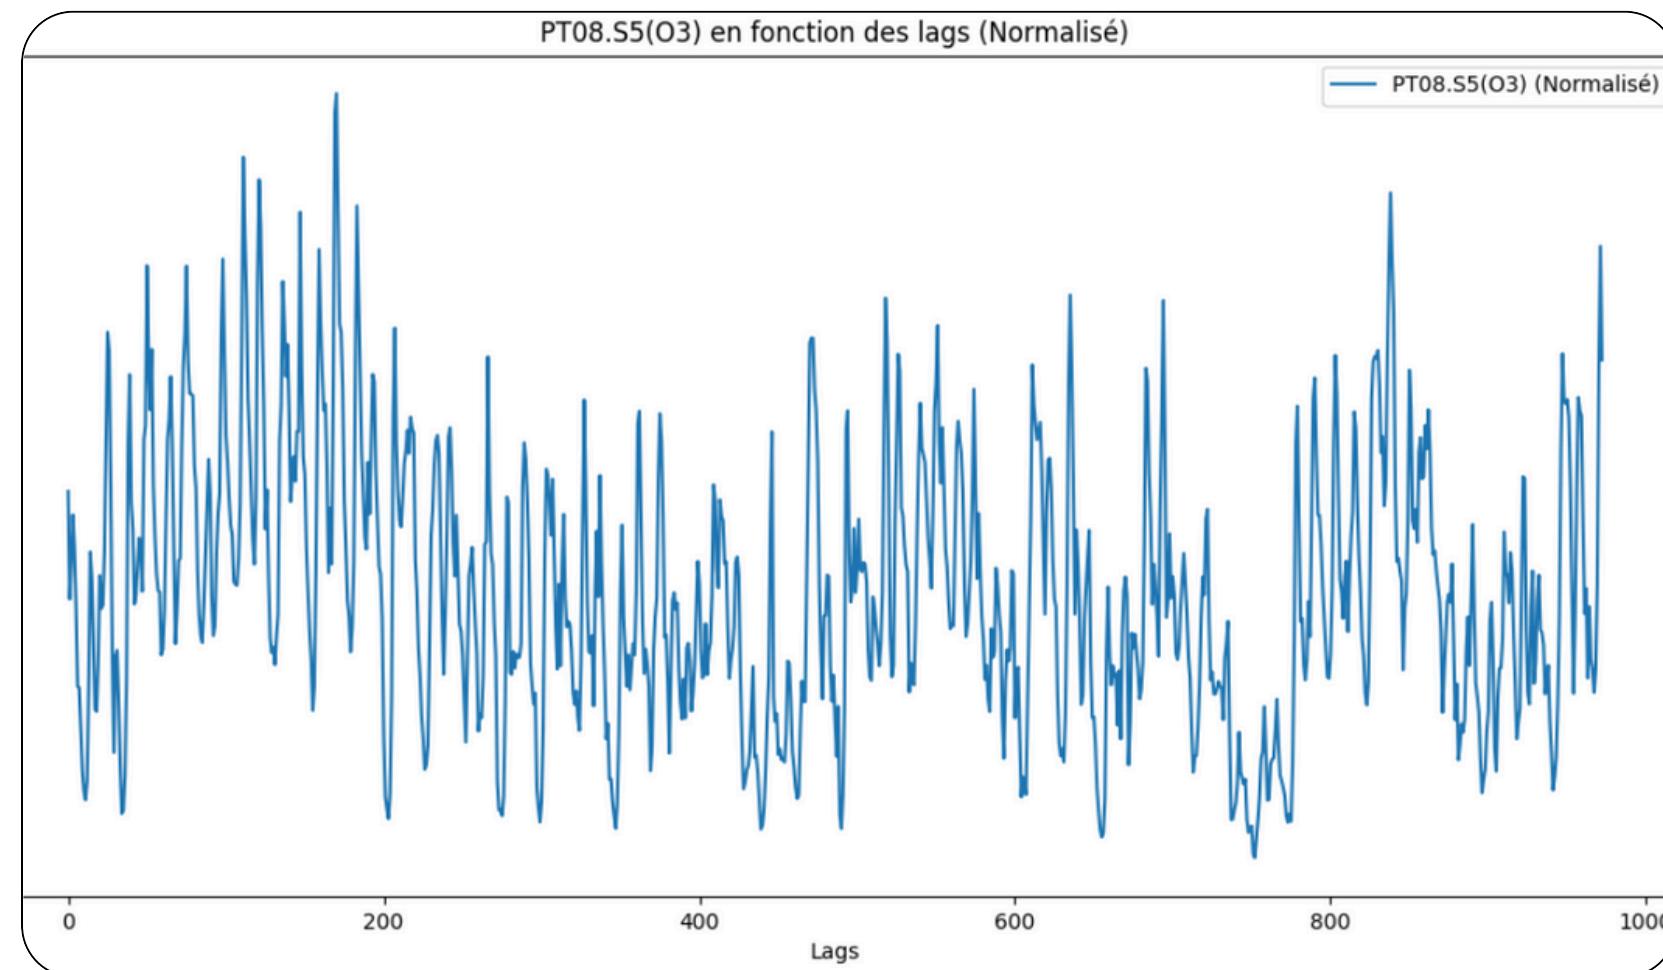
PT08.S5(03)	
Date	
2004-03-10	1267.50
2004-03-10	972.25
2004-03-10	1074.00
2004-03-10	1203.25
2004-03-10	1110.00
...	...
2005-04-04	1728.50
2005-04-04	1269.00
2005-04-04	1092.00
2005-04-04	769.75
2005-04-04	816.00

9357 rows × 1 columns

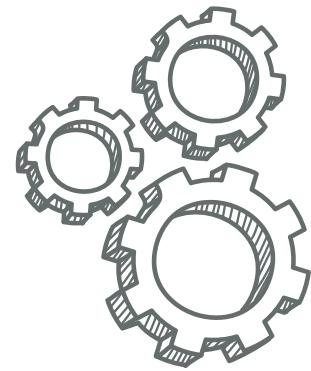


## PT08.S5(O3) en fonction des Lags

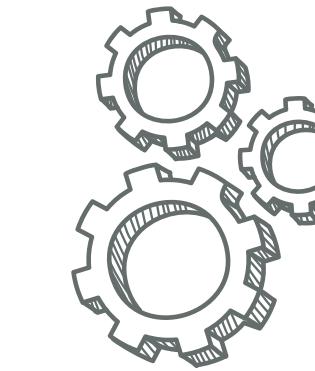
- Pics et creux réguliers
- Autocorrélation détectée



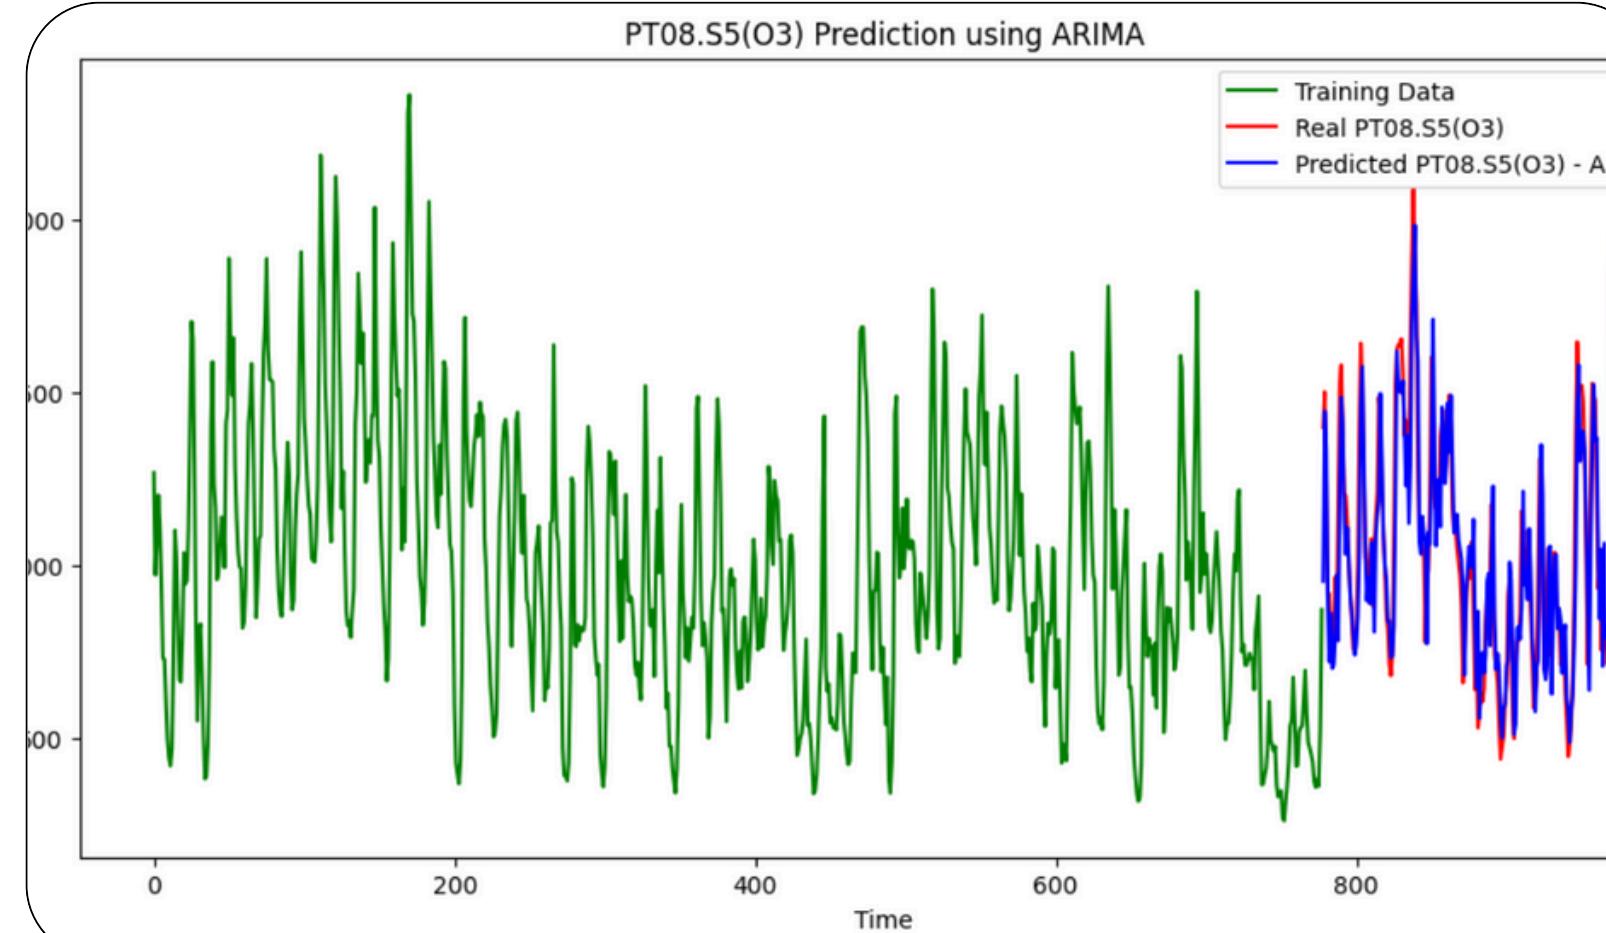
# Modélisation ARIMA



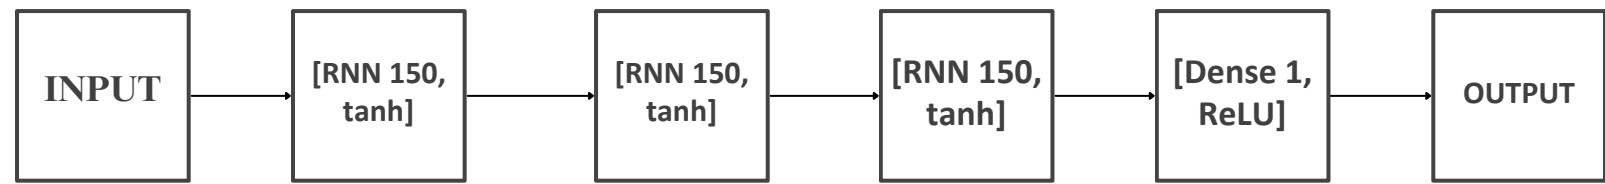
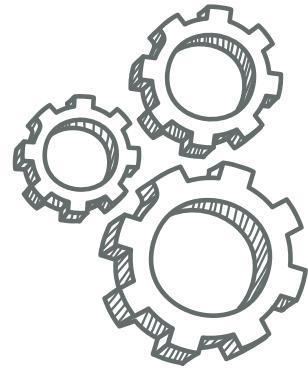
- **Architecture du modèle ARIMA :**
  - Modèle ARIMA (2,1,1) sélectionné via Grid Search en minimisant l'AIC .
- **Entraînement :**
  - Validation par glissement (Walk Forward Validation) .



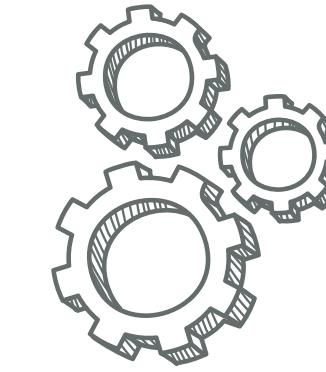
Residual Forecast Error: 16.099592752419454  
Forecast Bias: 16.099592752419454  
Mean Absolute Error: 123.4096448698548  
Mean Squared Error: 24965.311545484074  
Root Mean Squared Error: 158.00415040588038  
 $R^2$  Score: 0.7819570341404055



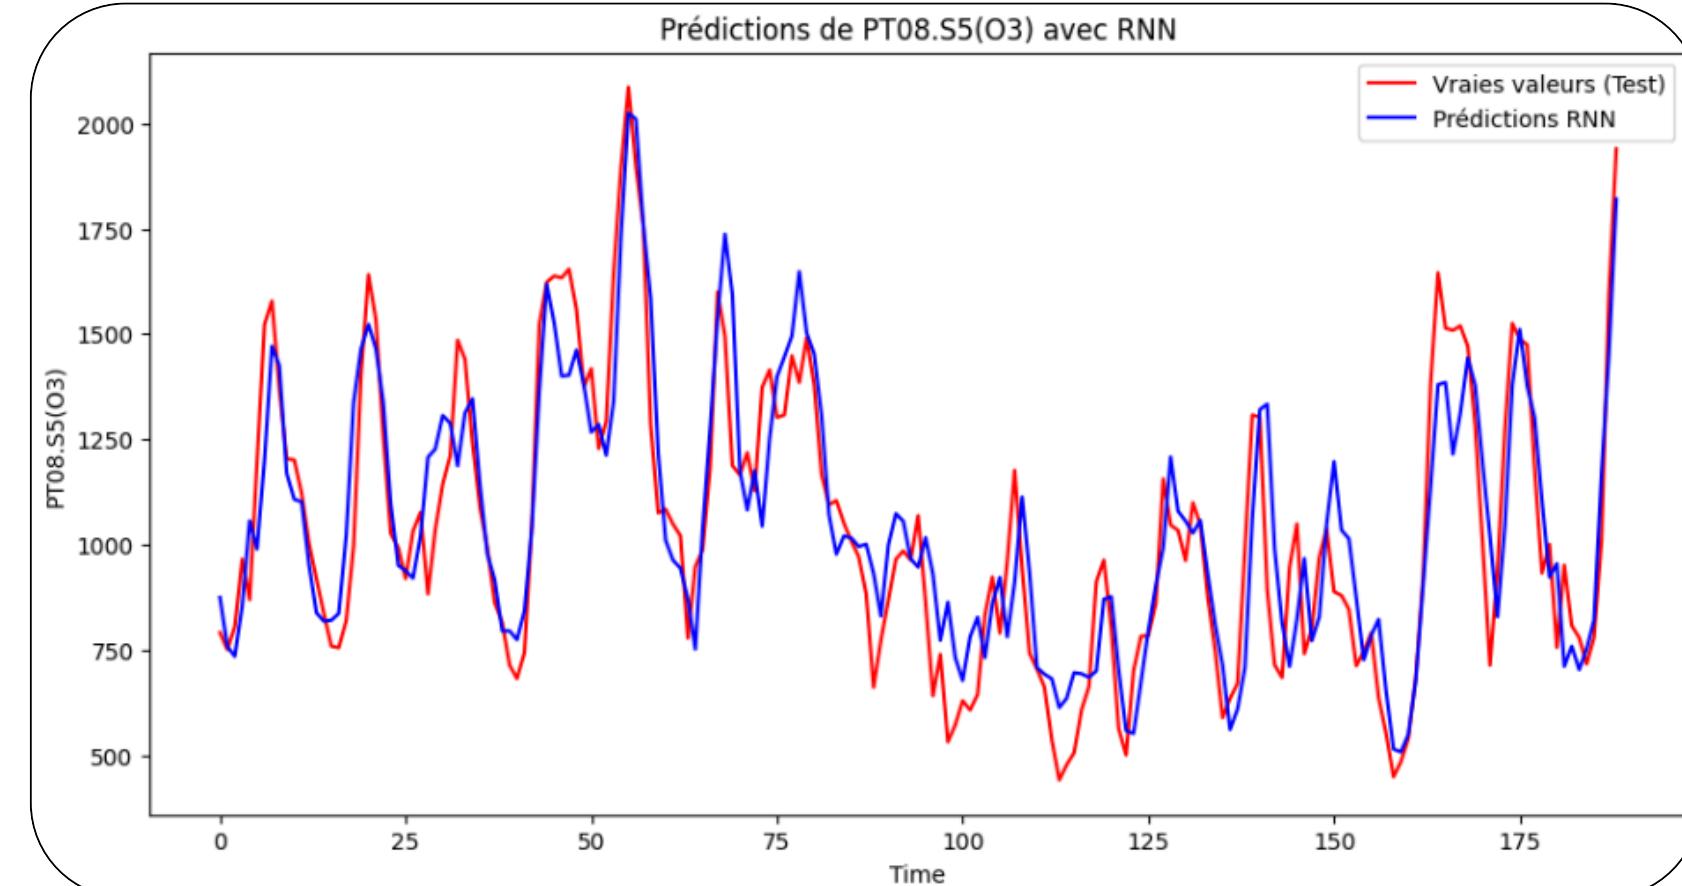
# Modélisation RNN



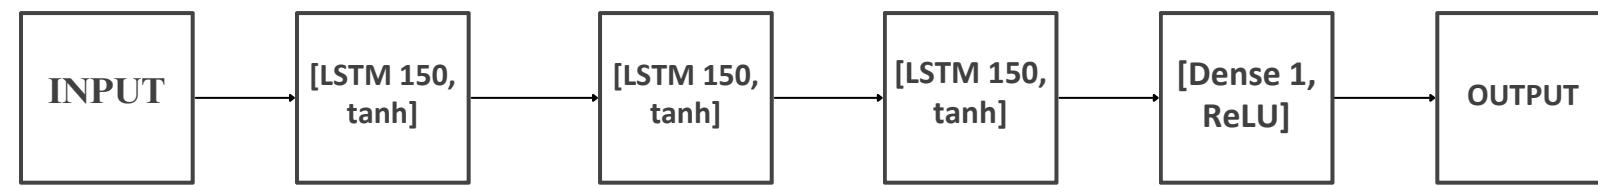
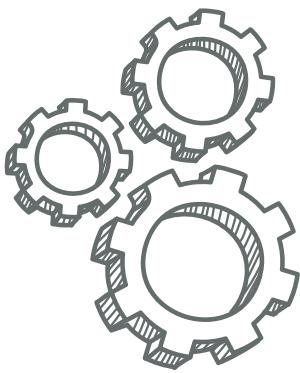
- Couches de Dropout (30%) pour éviter le surapprentissage.
- Optimisation avec Adam et fonction de perte MSE.
- 150 epochs, batch size=64, validation split=20%



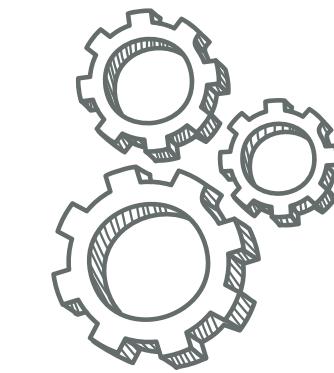
Evaluation du modèle:  
Mean Squared Error (MSE): 19888.14932535974  
Mean Absolute Error (MAE): 111.59860953326151  
 $R^2$  Score: 0.8460504789003471



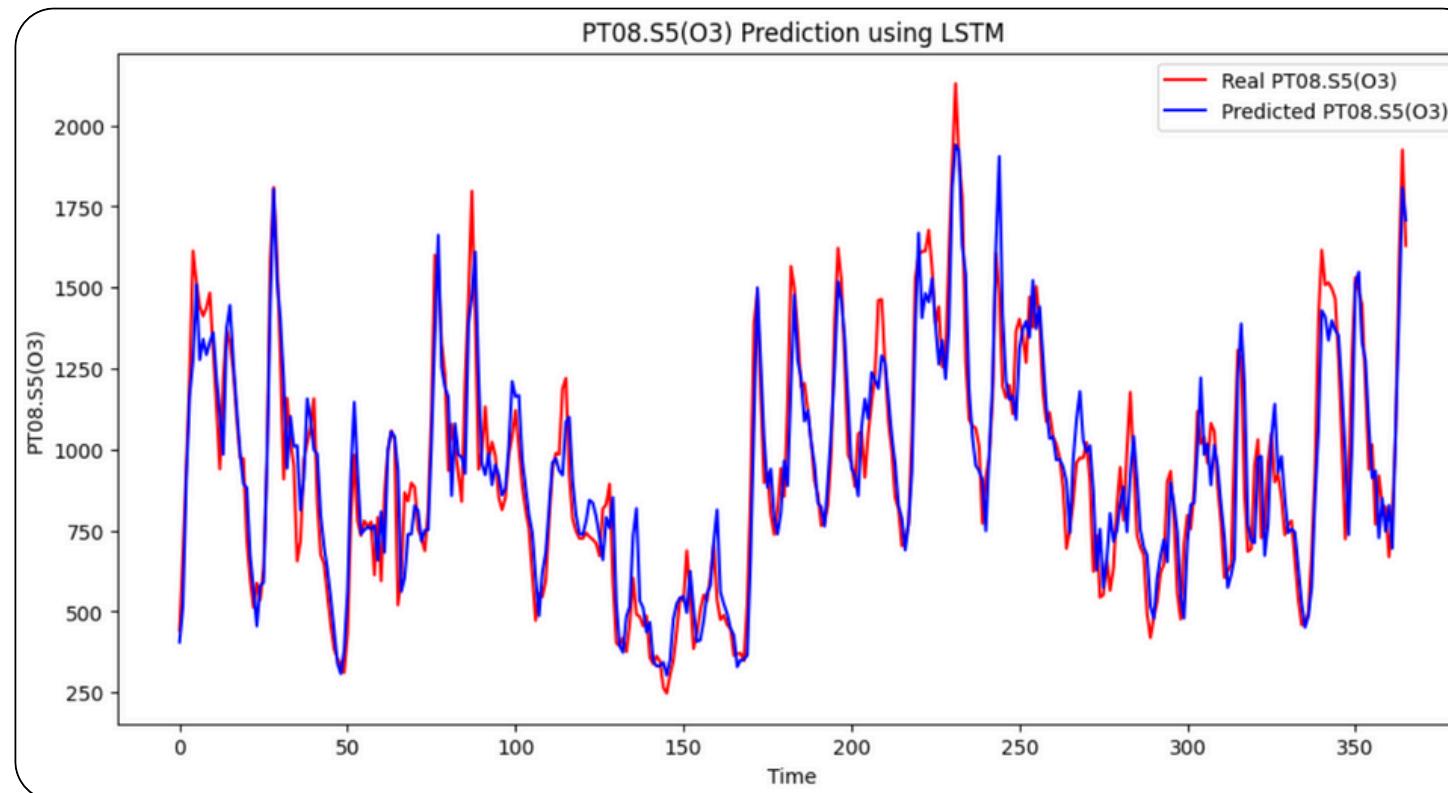
# Modélisation LSTM



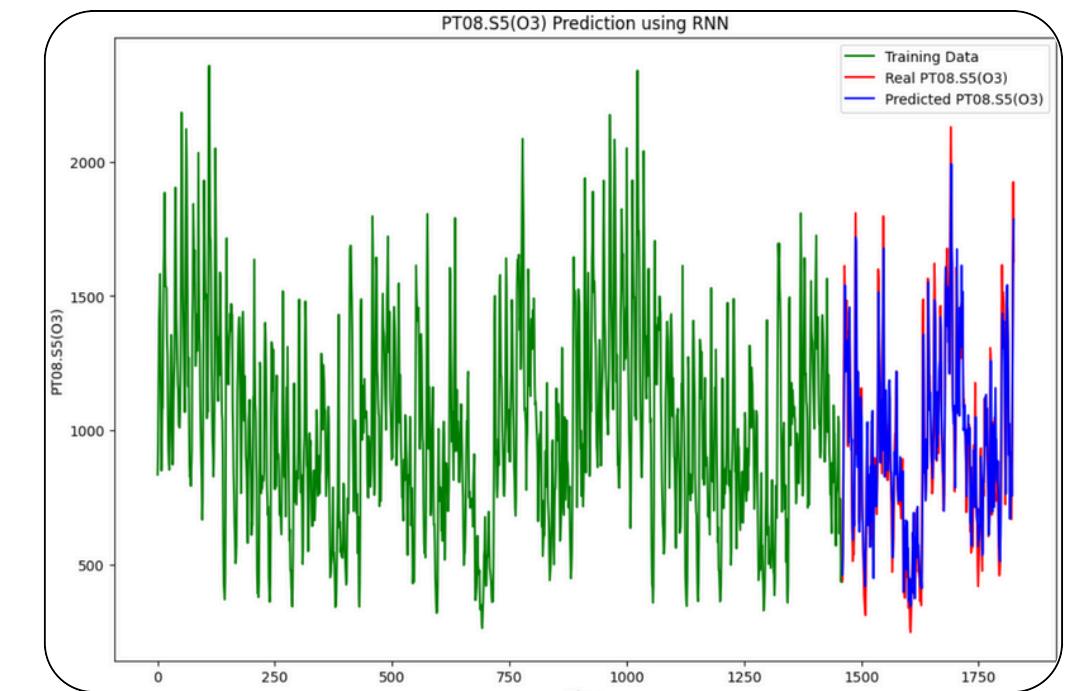
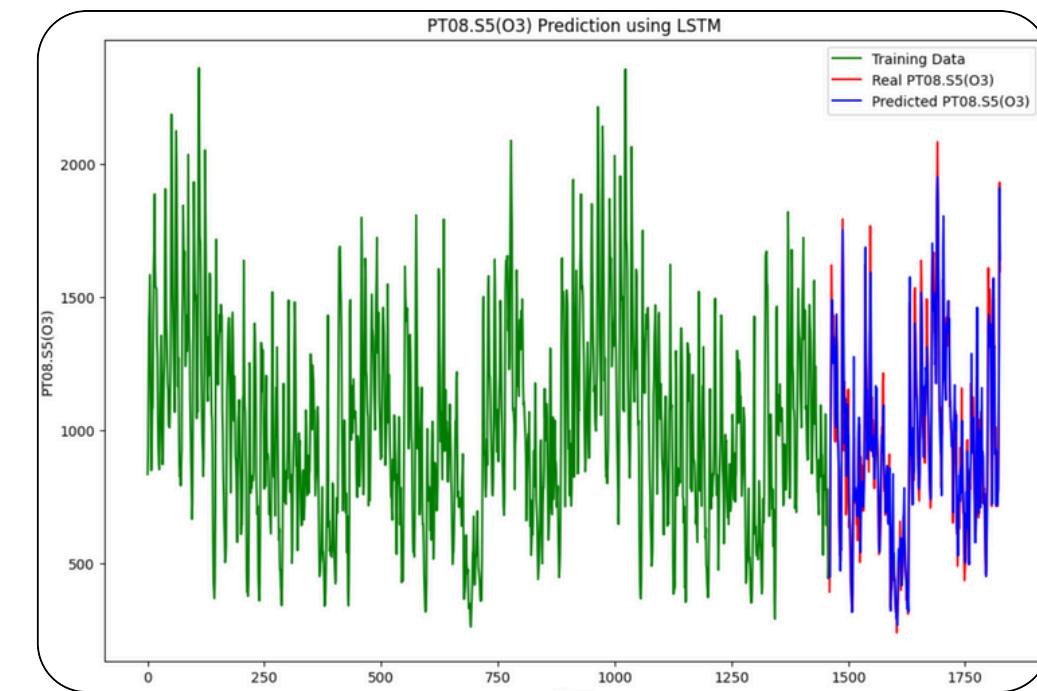
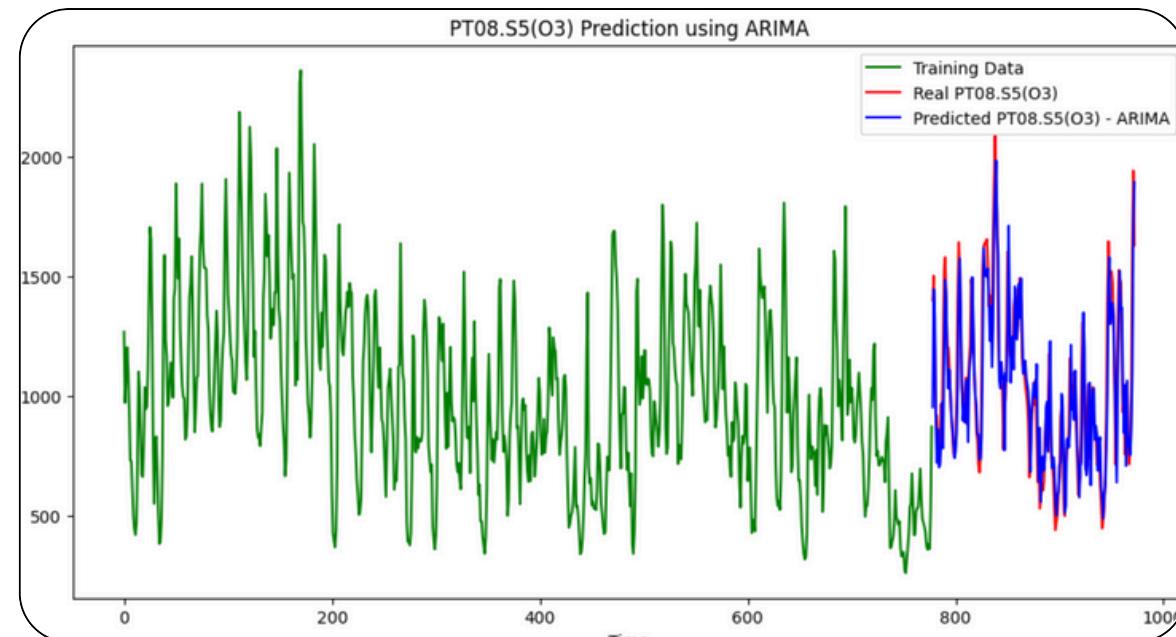
- Couches de Dropout (30%) pour éviter le surapprentissage.
- Optimisation avec Adam et fonction de perte MSE.
- 150 epochs, batch size=64, validation split=20%



Evaluation du modèle:  
Mean Squared Error (MSE): 17572.524961459403  
Mean Absolute Error (MAE): 102.97603910010308  
 $R^2$  Score: 0.8639751865258364



# Comparaison finale des performances



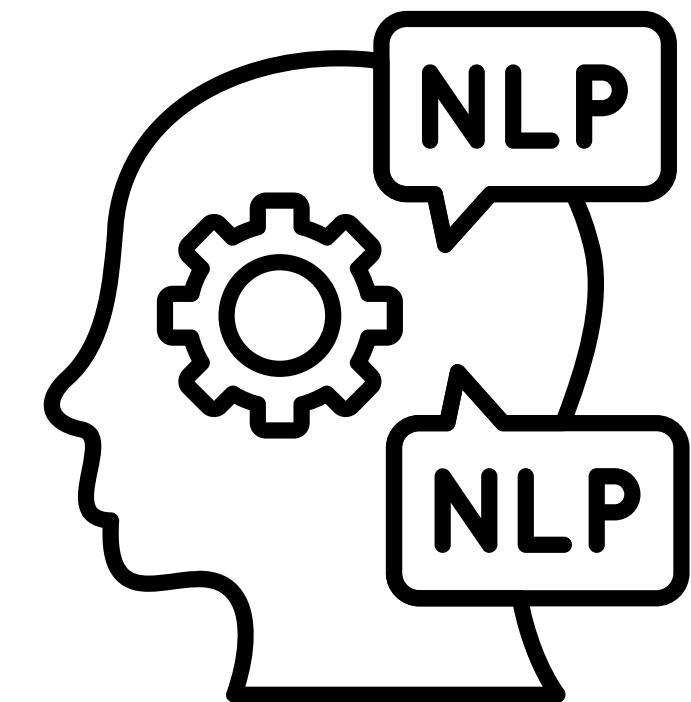
ARIMA

LSTM

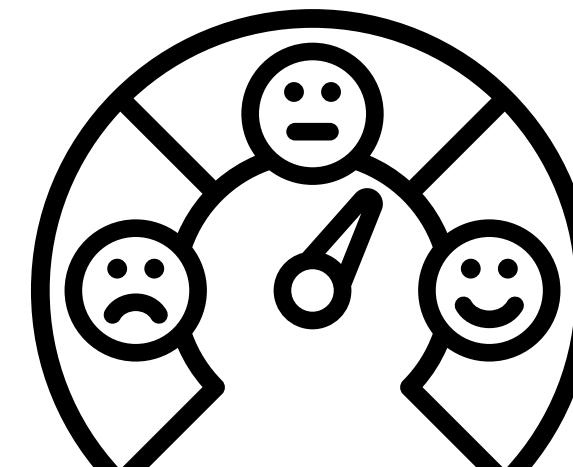
RNN

Model	MAE	MSE	RMSE	R2
ARIMA	123.409645	24965.311545	158.004150	0.781957
RNN	115.913238	21282.121348	145.883931	0.835376
LSTM	102.976039	17572.524961	132.561401	0.863975

## 6.NLP et Techniques Avancées



# Sentiment Analysis



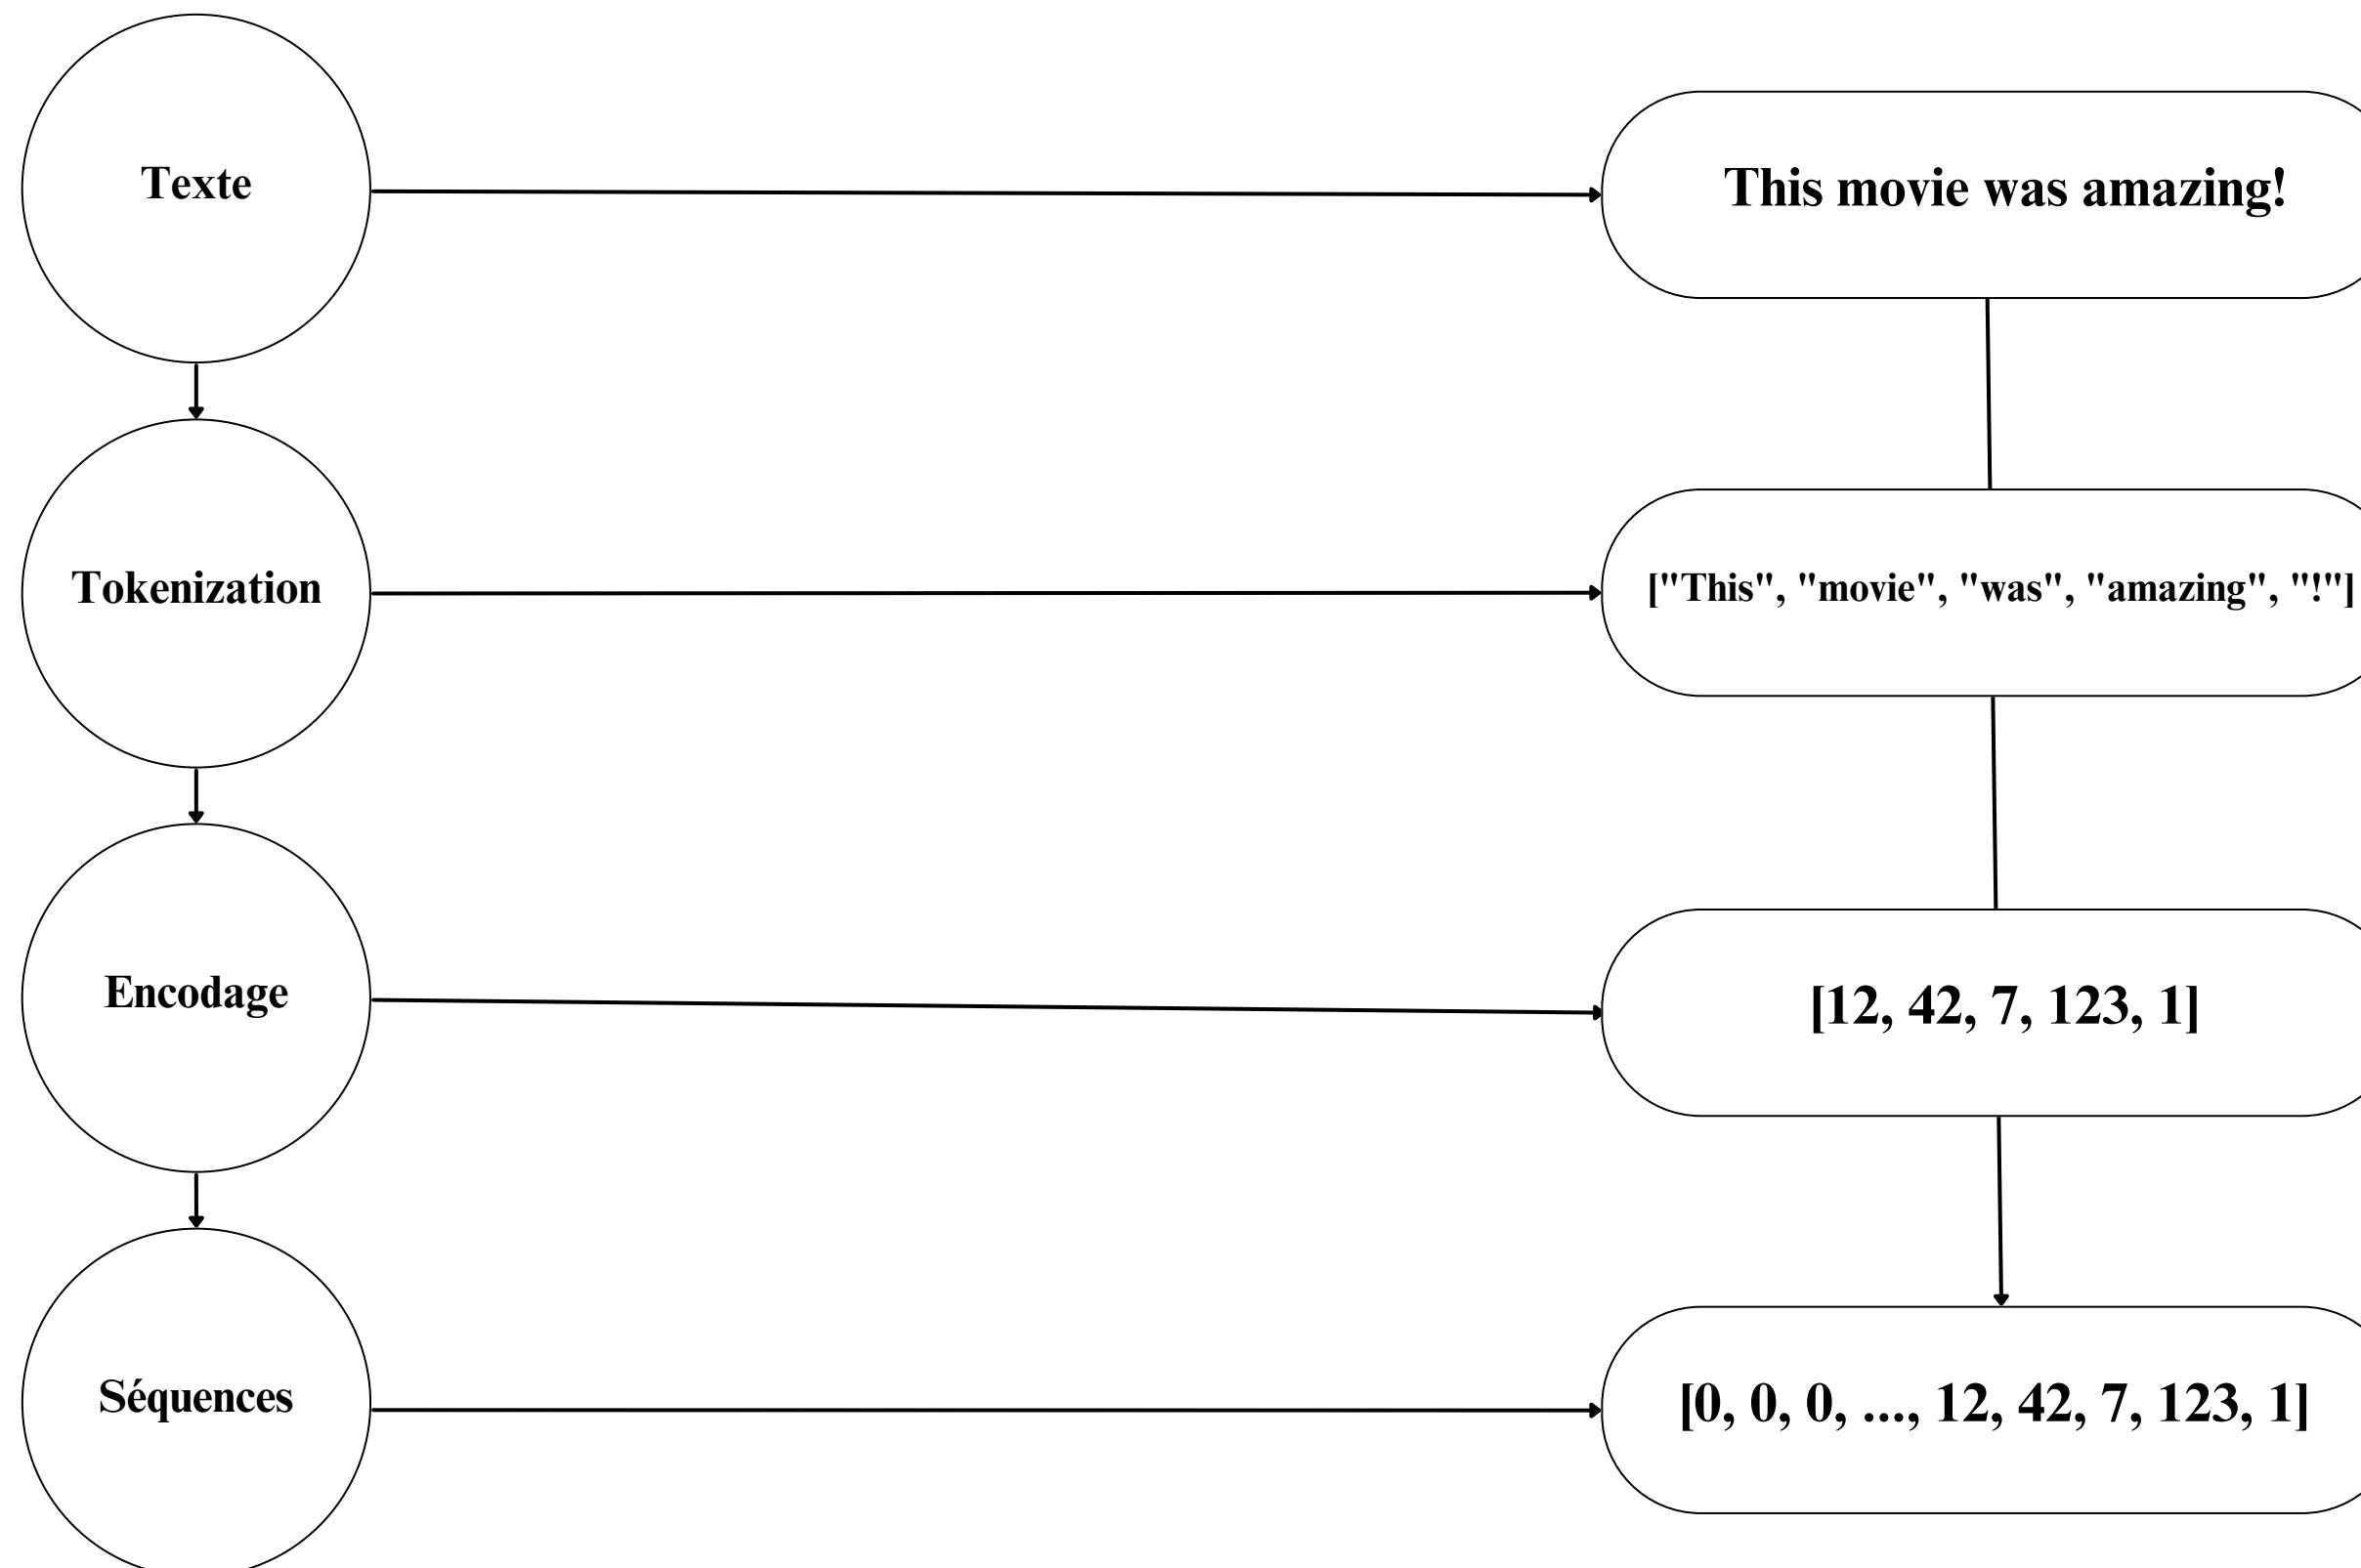
## Objectif



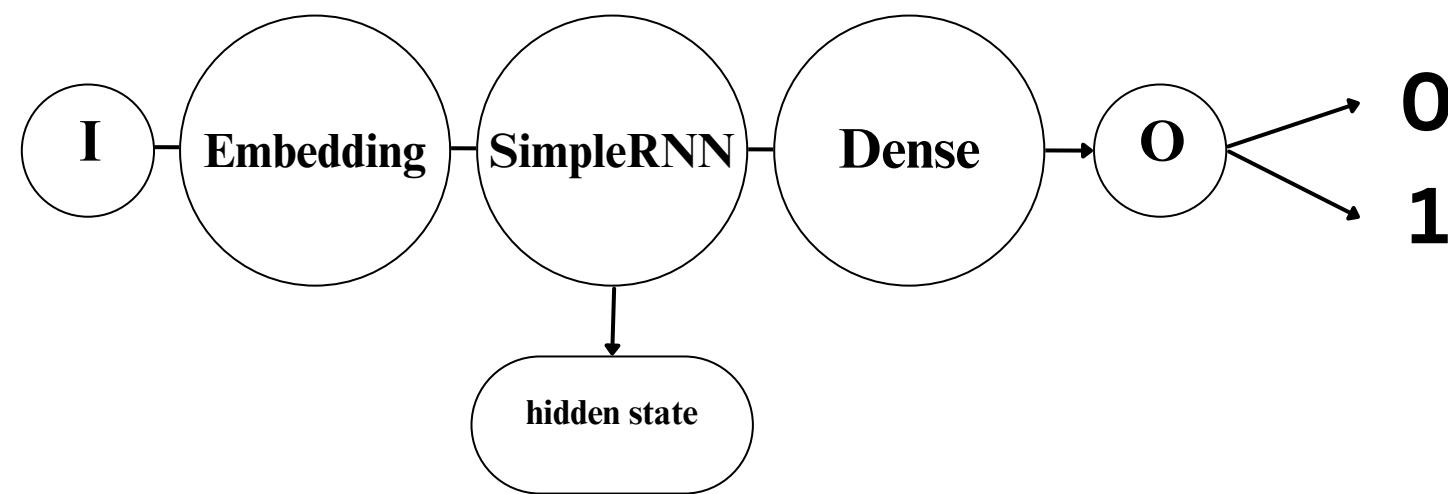
- Analyser les sentiments des critiques de films (positif/négatif).
- Comparer les performances des modèles RNN, LSTM et GRU.
- Intégrer un mécanisme d'attention pour améliorer les résultats.
- Utilisation de transformers .



## Préparation et Transformation



## Modélisation avec RNN et LSTM

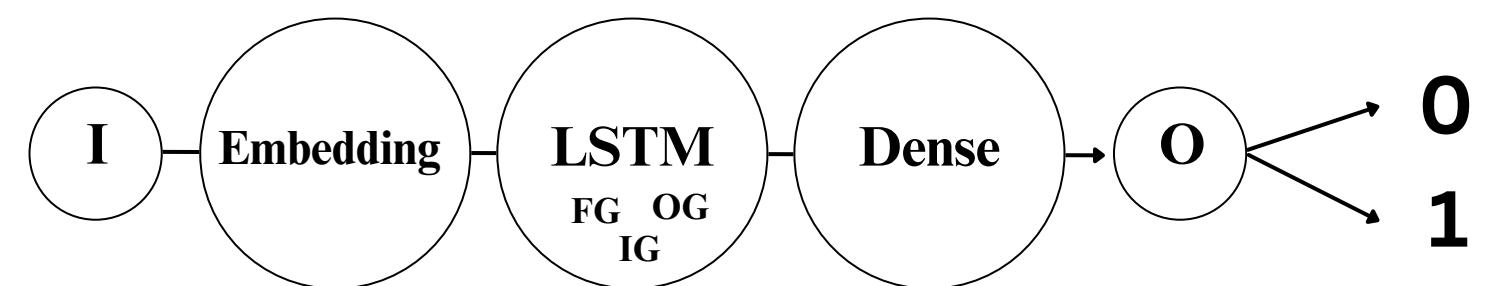


```
positive_review = "it was amazingly very great"
predict(positive_review)

negative_review = "that movie really sucked. I hated it and wouldn't watch it again."
predict(negative_review)
```

```
1/1 [0.9949482]
  
1/1 [0.06237121]

```



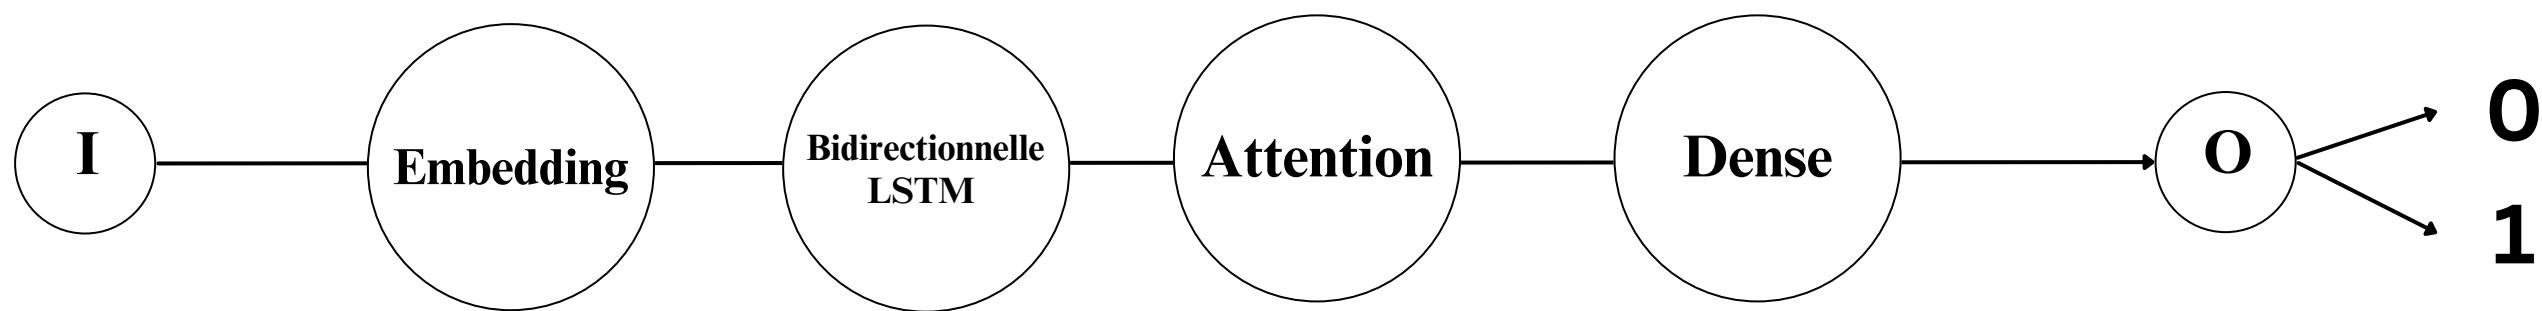
```
positive_review = "That movie was! really loved it and would great watch it again because it was ama
predict(positive_review)

negative_review = "that movie really sucked. I hated it and wouldn't watch it again. Was one of the v
predict(negative_review)
```

```
1/1 [0.9041555]
  
1/1 [0.03734026]

```

## Mécanisme d'Attention



```
positive_review = "That movie was! really loved it and would great watch it again"
predict(model_lstm_attention,positive_review)
```

```
negative_review = "that movie really sucked. I hated it and wouldn't watch it again"
predict(model_lstm_attention,negative_review)
```

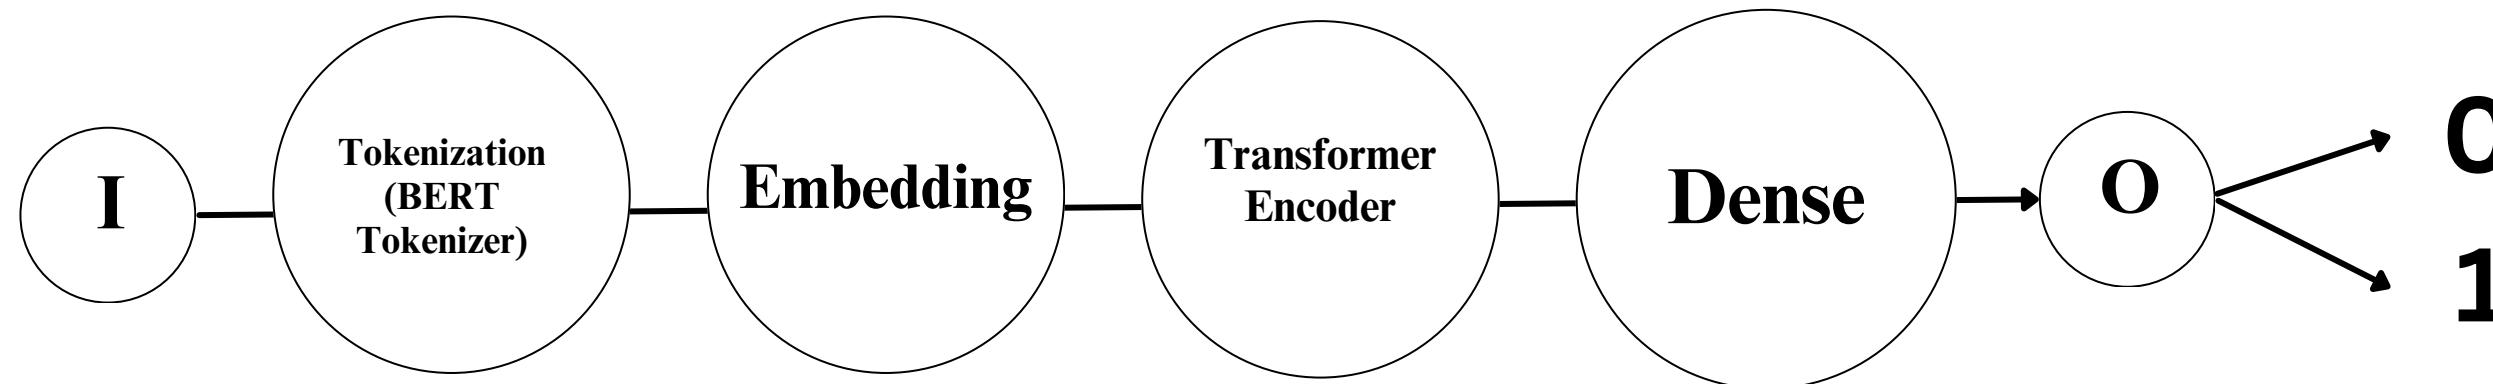
```
→ 1/1 ━━━━━━ 0s 64ms/step
[0.9814324]
```



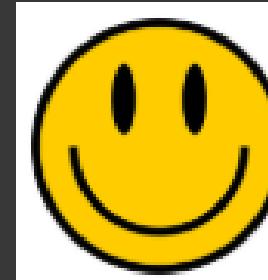
```
1/1 ━━━━━━ 0s 62ms/step
[0.04269176]
```



## Modélisation avec Transformers



```
# Exemple de prédiction
text = "the thing that really ? people and this film proves that as the direc
print(predict_sentiment(text))
text = "This movie was absolutely bad! I hated every minute of it."
print(predict_sentiment(text))
```

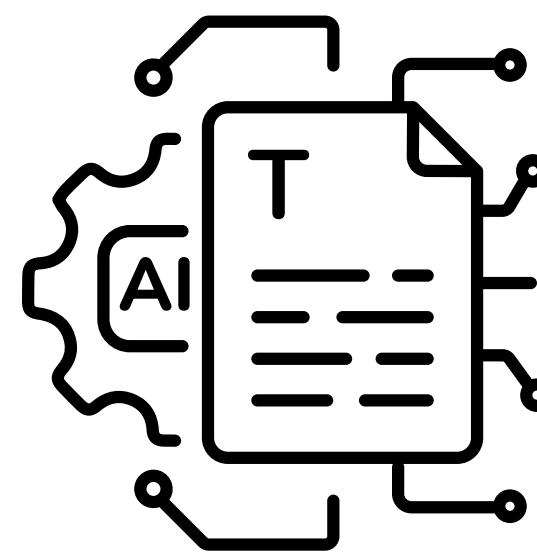


('Positive', <tf.Tensor: shape=(), dtype=float32, numpy=0.801487922668457>)

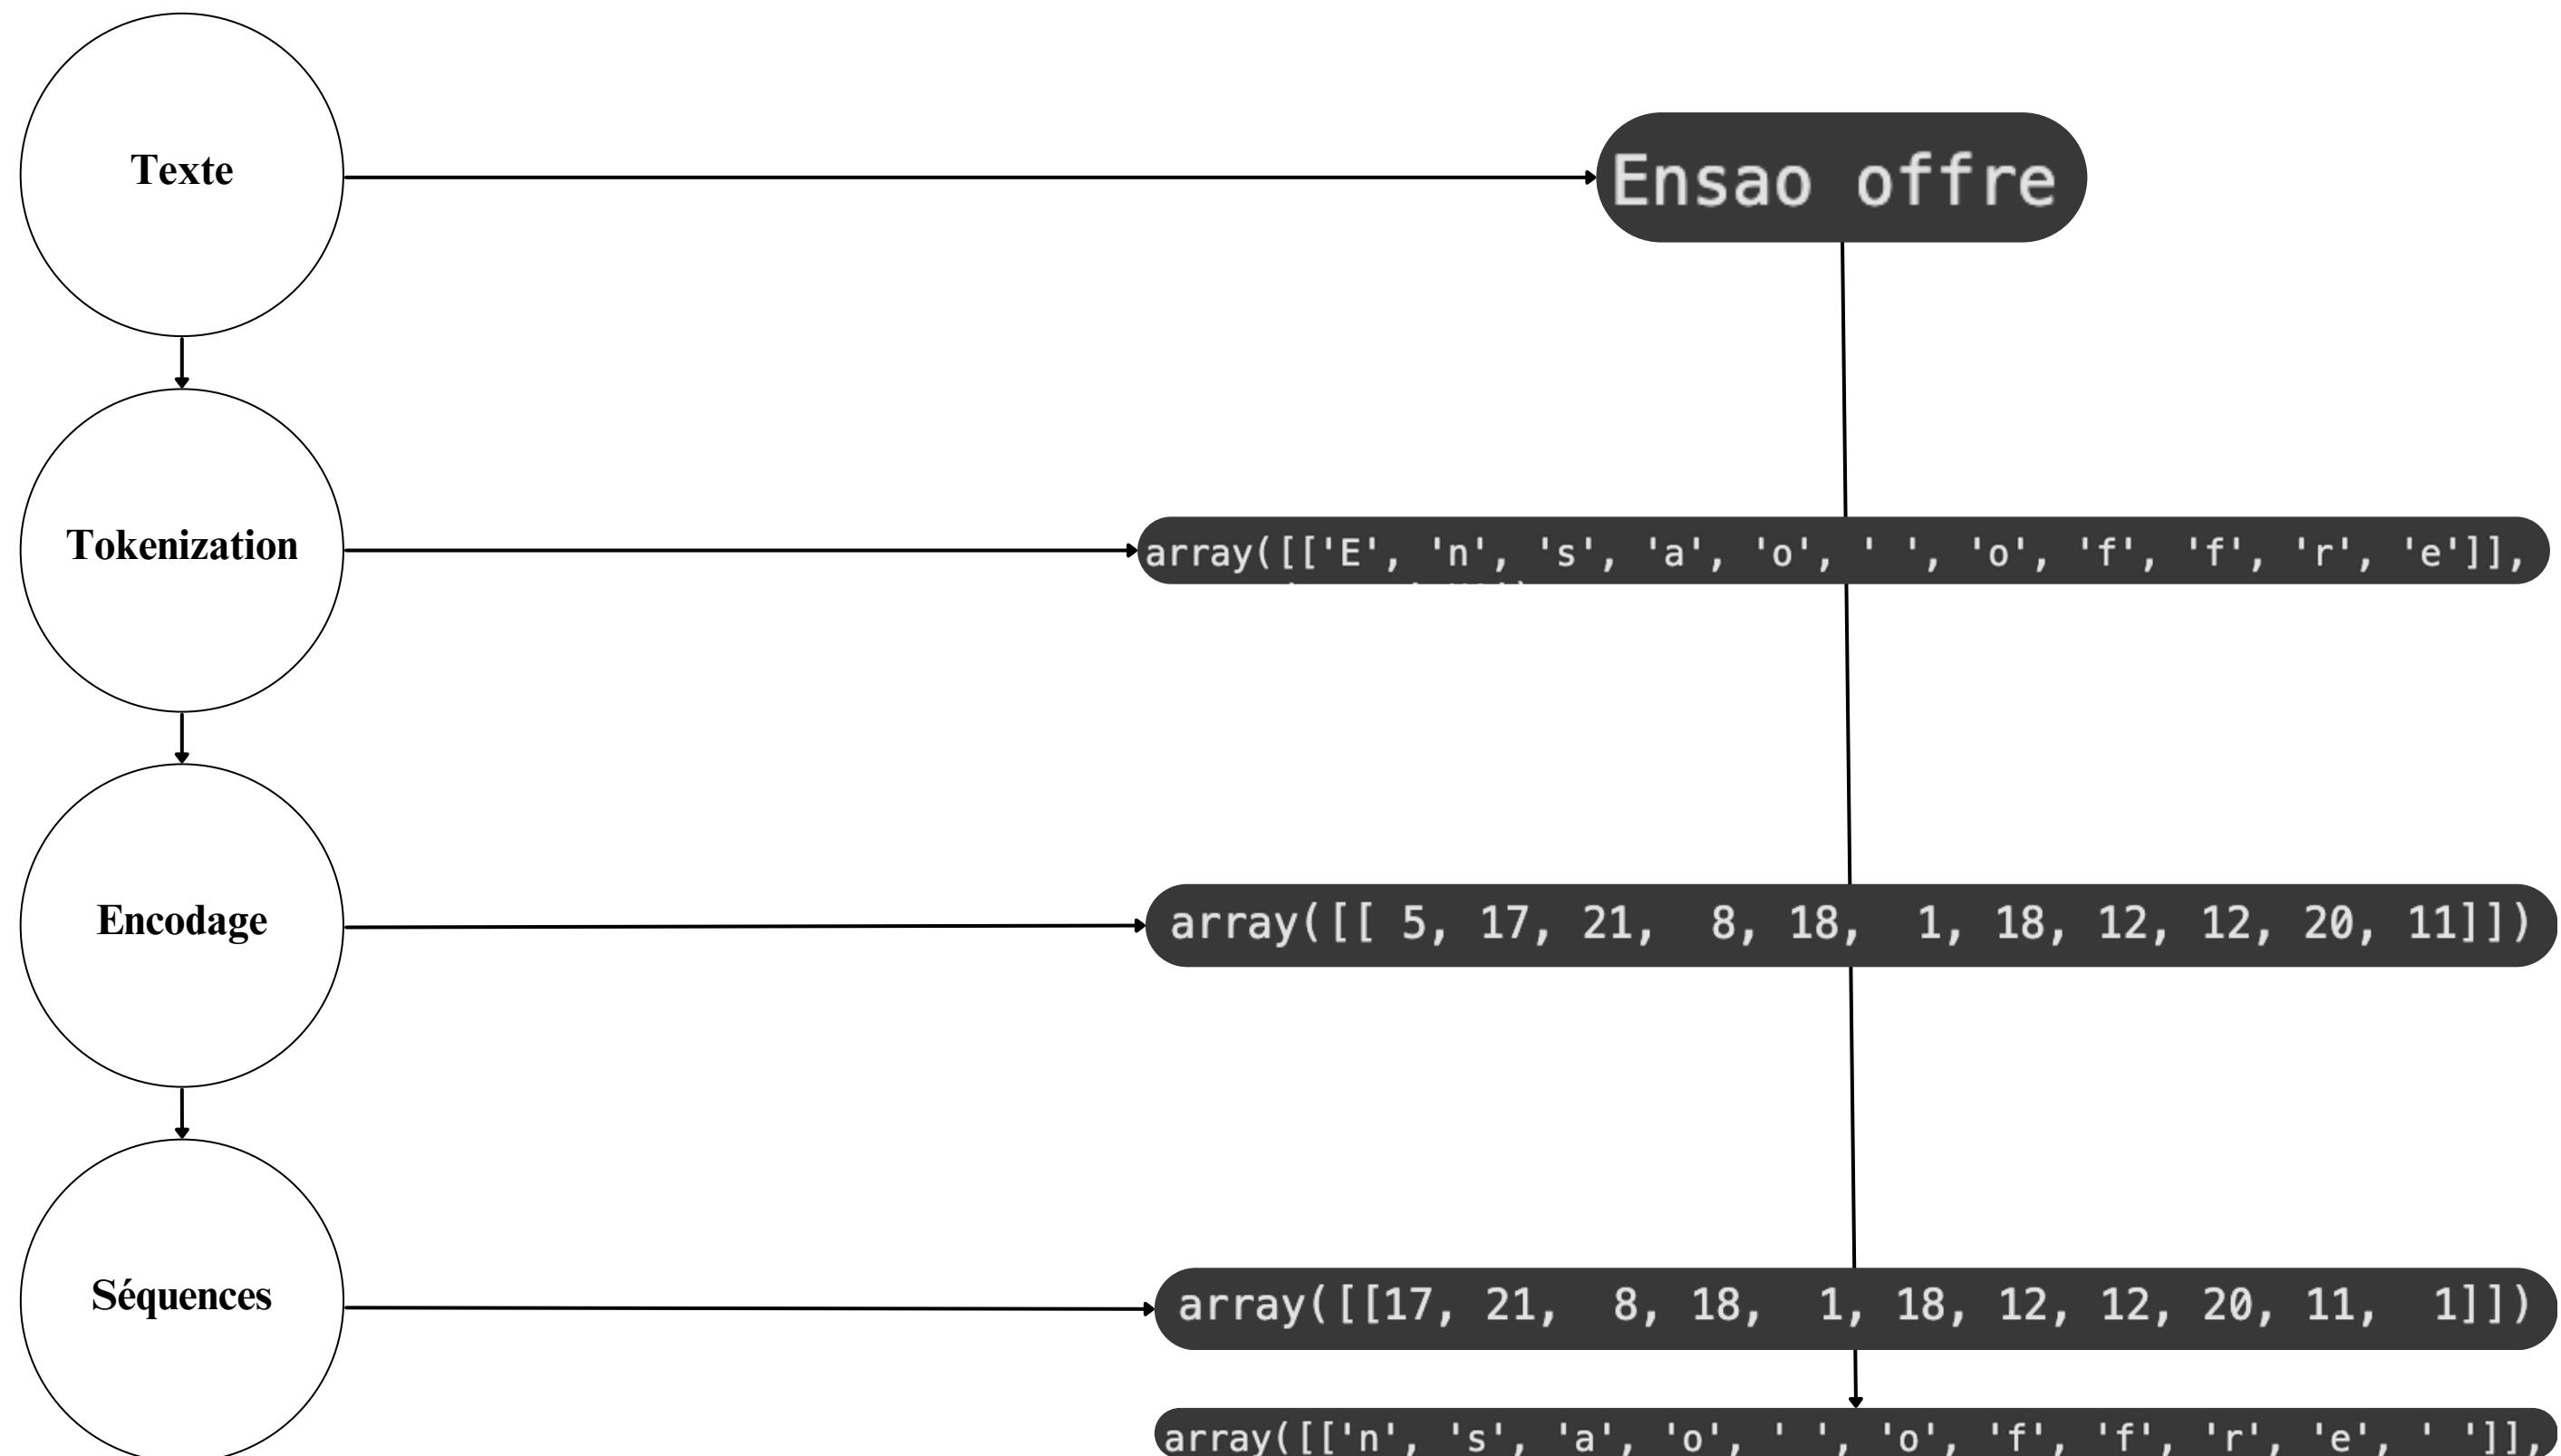


('Negative', <tf.Tensor: shape=(), dtype=float32, numpy=0.23638194799423218>)

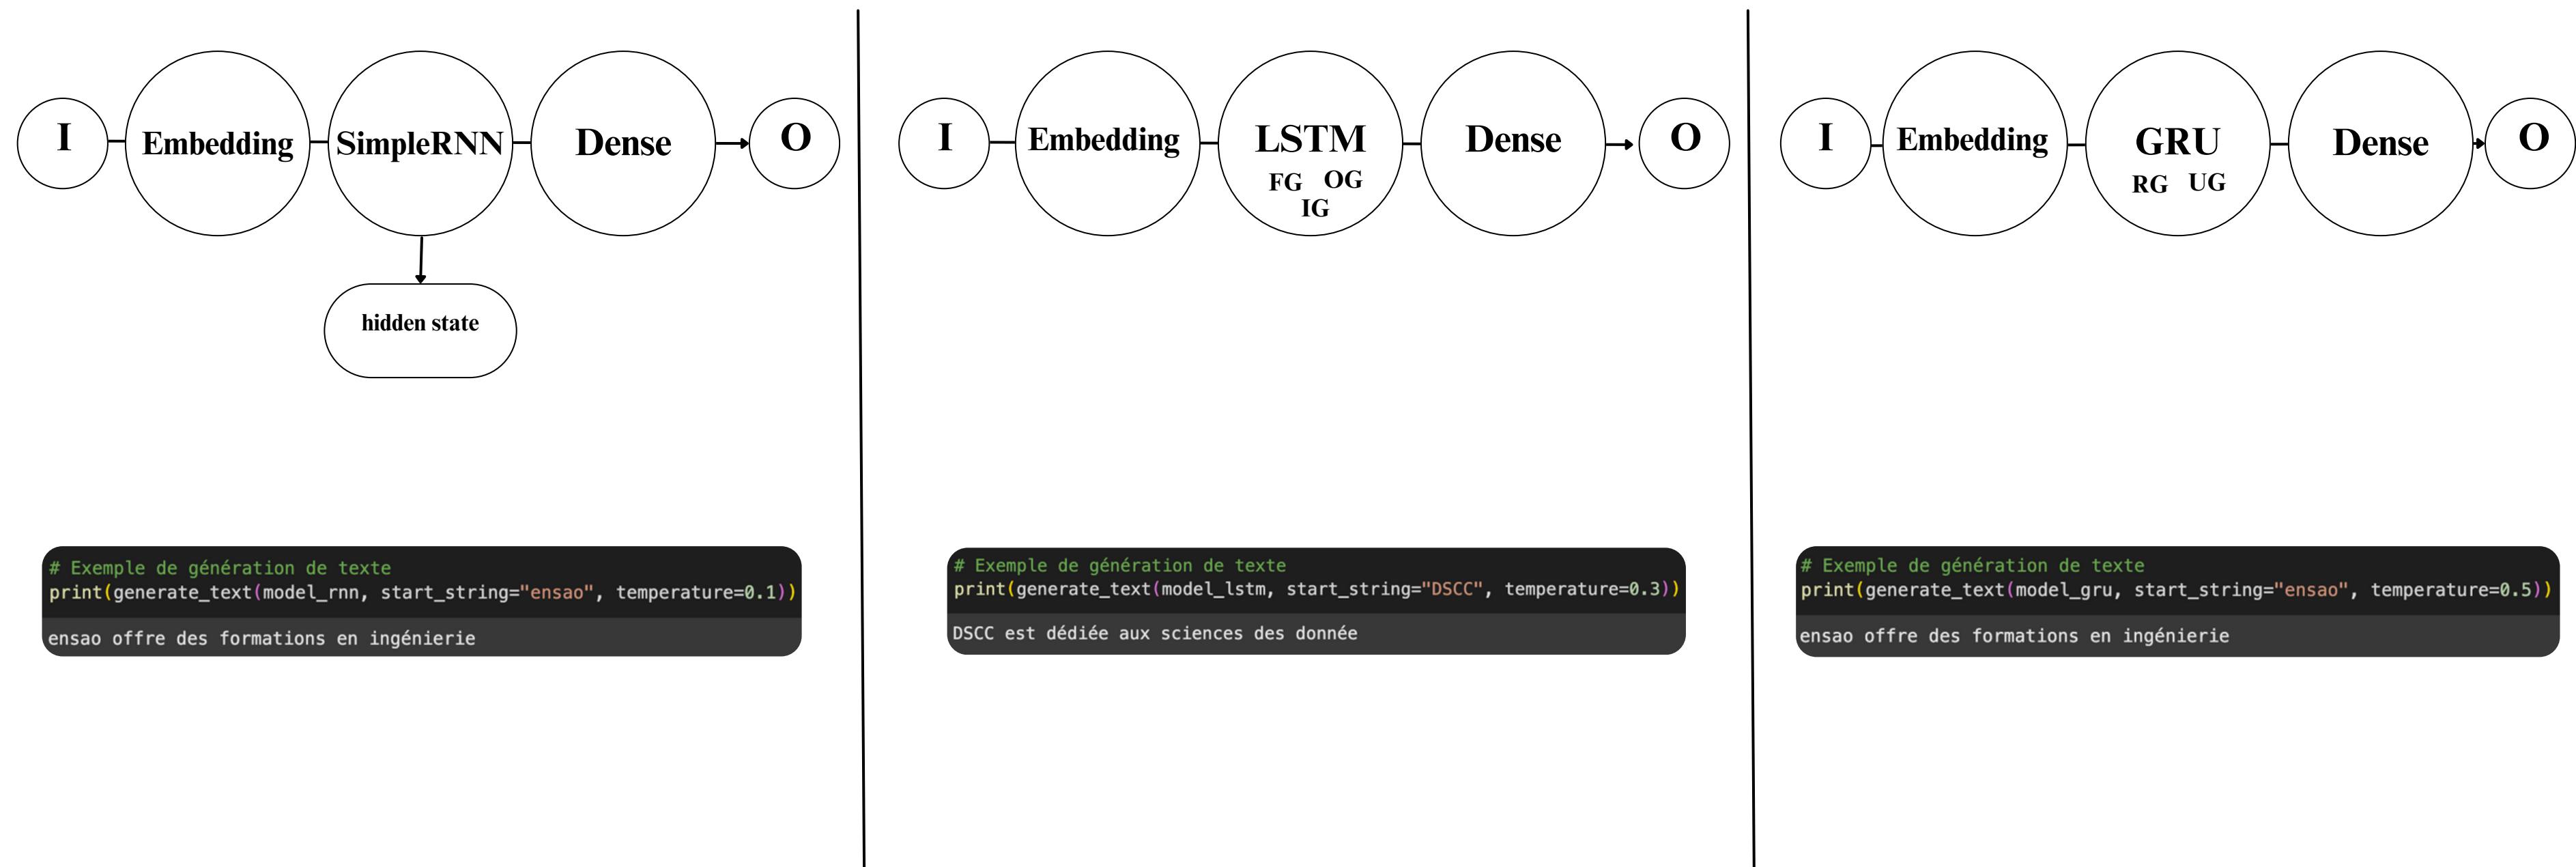
# **Text Generation**



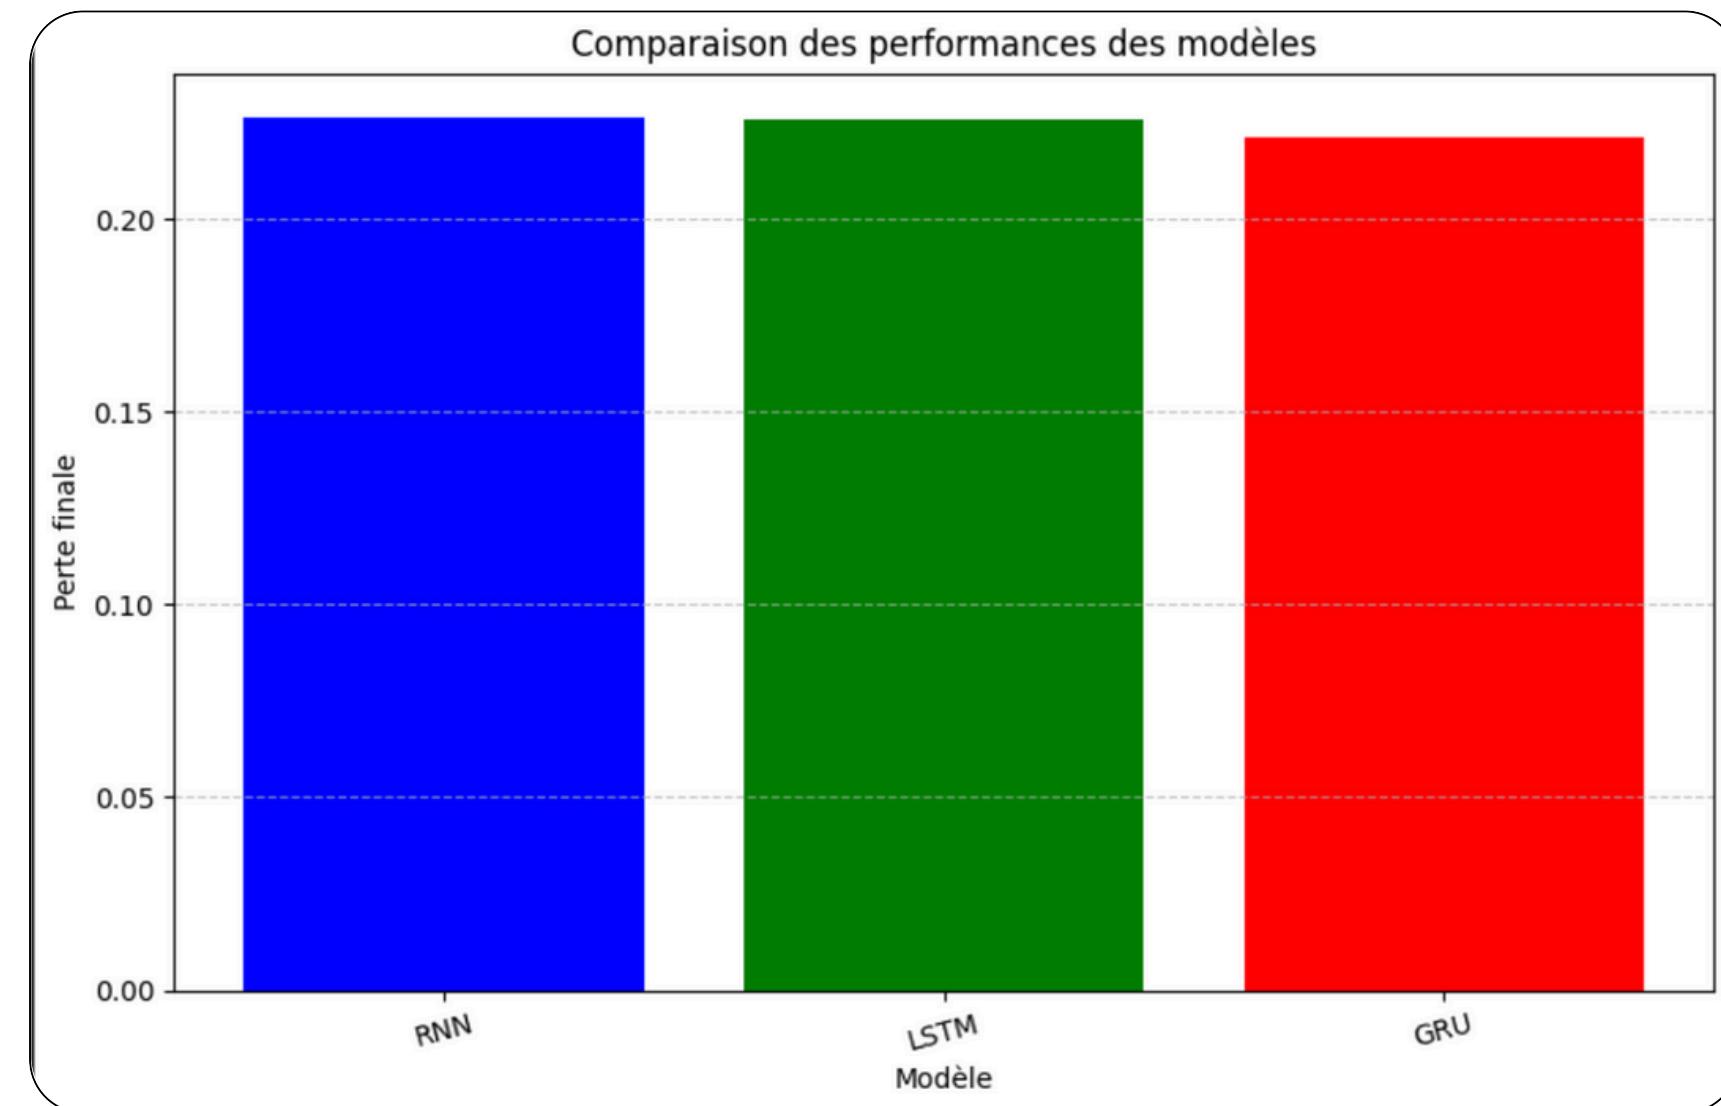
## Préparation et Transformation



# Modélisation avec RNN, LSTM et GRU

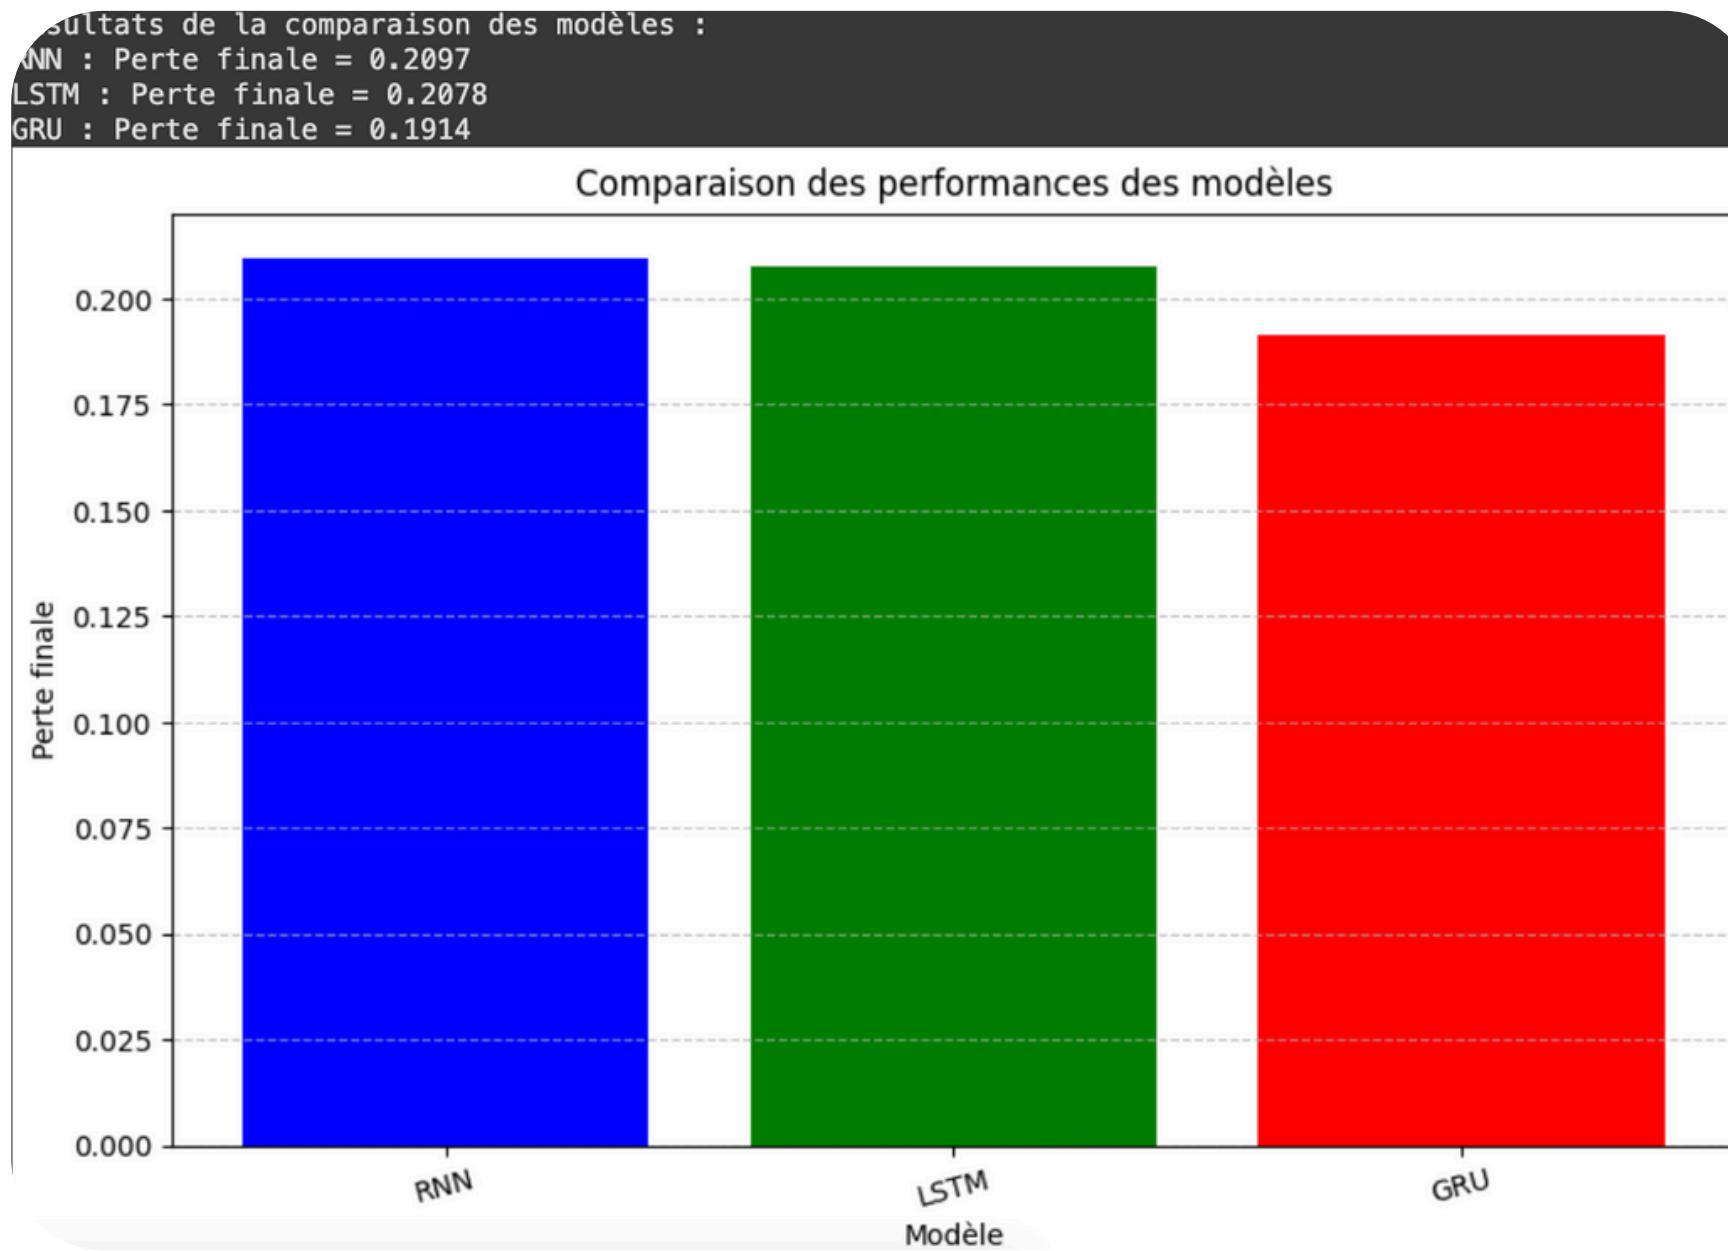


## Comparaison des résultats

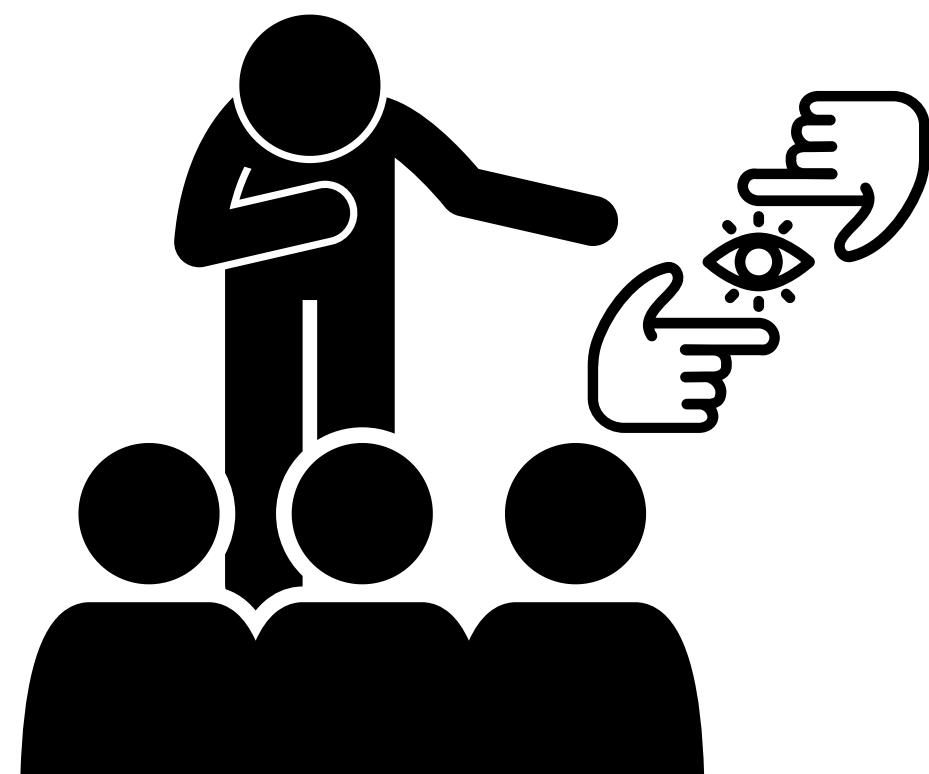


Résultats de la comparaison des modèles :  
RNN : Perte finale = 0.2261  
LSTM : Perte finale = 0.2226  
GRU : Perte finale = 0.2215

## Utilisation de GloVe



## 7. Perspectives et Conclusions



**Limites**

**Vanishing Gradient**

**Complexité**

**Interprétabilité**

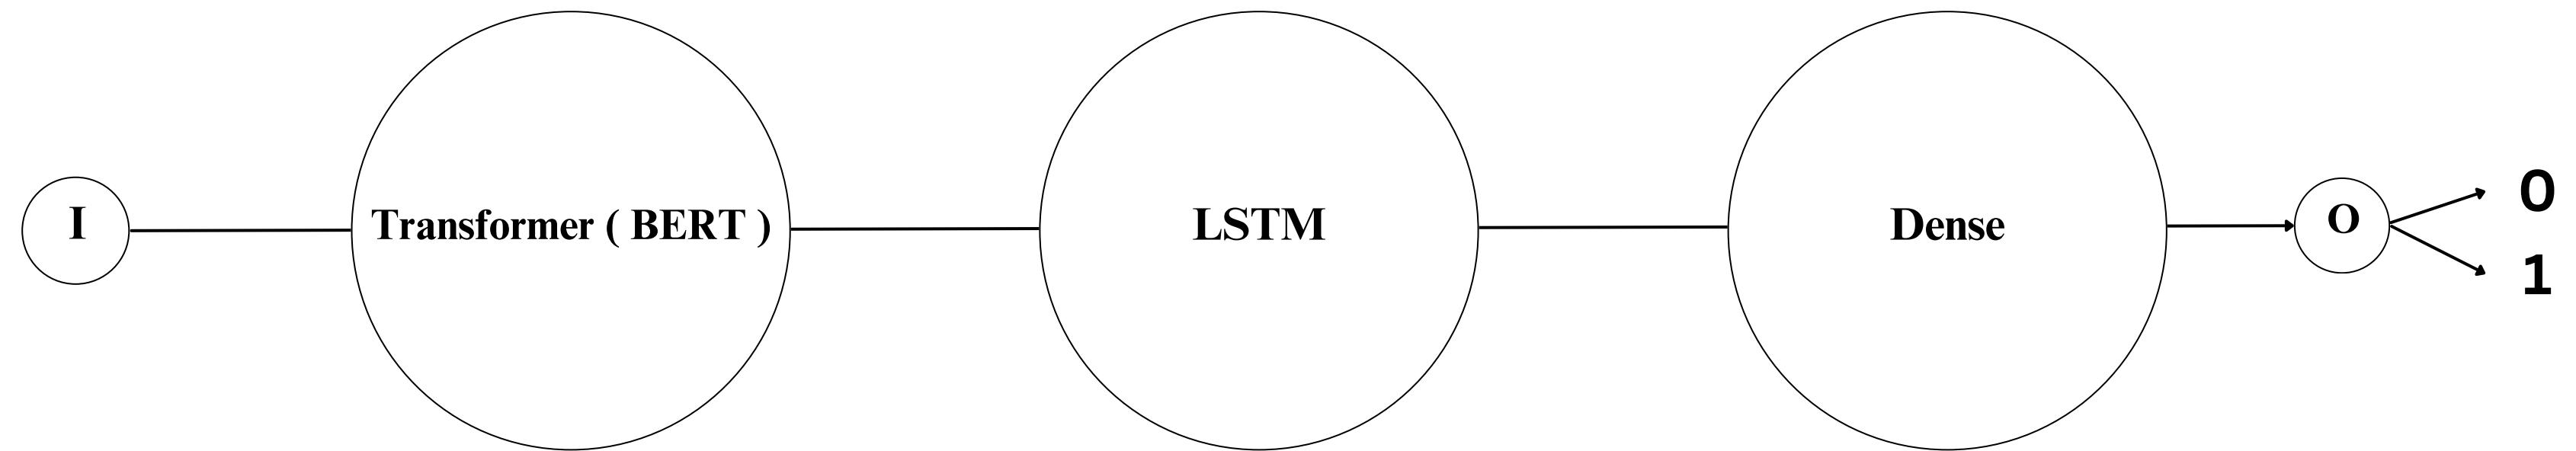
**Sensibilité**

## **Modèles Préentraînés**

**Transfer Learning**

**Contexte Global**

## **Transformers Hybrides**



MERCI POUR VOTRE ATTENTION

